



Compiling and running OpenMP codes on Kebnekaise at HPC2N

Pedro Ojeda & Joachim Hein

High Performance Computing Center North &
Lund University



Overview

- HPC2N's Kebnekaise system
- The module environment
- Building OpenMP executables on HPC2N systems
- Executing OpenMP jobs on HPC2N system



Kebnekaise@ HPC2N

- 432 compute nodes
 - 52 compute-skylake nodes
 - 20 large-memory nodes
 - 32 2xK80 GPU nodes
 - 4 4xK80 GPU nodes
 - 10 2xV100 GPU nodes
 - 36 KNL nodes
-
- Compute nodes (Broadwell)
 - Intel Xeon E5-2690 v4 processors
 - 14 cores per processor (2 sockets)
 - 28 cores per node
 - 128 GB memory per node



Connecting to Kebnekaise: Thinlinc client

- Download and install the client from:
<http://www.cendio.com/downloads/clients/>
- Launch the client
- Enter **kebnekaise-tl.hpc2n.umu.se** in the server field.
- Just for the first time, go to “Options” button in the client dialog box, then “Screen” and uncheck “Full screen mode”
- Enter your user-id & password
- Click [Connect]

Detailed descriptions:

<https://www.hpc2n.umu.se/documentation/guides/thinlinc>



Connecting to Kebnekaise: Command prompt

- Kebnekaise is a Linux based machine
- Connect to it via: ssh
- Address: **kebnekaise.hpc2n.umu.se**
- Linux/Mac OSX: run `ssh -Y username` in terminal window
- Windows: install putty to run `ssh -Y username`
- Correct password
- Detailed description:
 - <https://www.hpc2n.umu.se/documentation/guides/mac-connection> (Mac)
 - <https://www.hpc2n.umu.se/documentation/guides/windows-connection> (Windows)
 - <https://www.hpc2n.umu.se/documentation/guides/linux-connection> (LINUX)



Building OpenMP executables

- Use standard compilers (e.g. Intel, Gnu, ...) to compile the source
- Enable OpenMP via compiler flag
- HPC2N systems use modules to make these steps convenient



Hands on: compiling OpenMP code

- Load the compiler

Compiler	Kebnekaise example
GNU 10.2.0	module load foss/2020b
intel/2020b	module load intel/2020b

- Examples for compiling code

Case	Command
C, GCC compiler	gcc -O3 -march=native -fopenmp -o prog prog.f90
F90, GNU compiler	gfortran -O3 -march=native -fopenmp -o prog prog.f90
C, Intel compiler	icc -qopenmp -O3 -xHost -o prog prog.c
F90, Intel compiler	ifort -qopenmp -O3 -xHost -o prog prog.f90

Rem: specify other flags as usual, e.g. optimisation



OpenMP and CMake

- Good OpenMP support for C, C++ and Fortran in CMake
- The cmake module FindOpenMP sets the variables
 - OpenMP_C_FLAGS - flags to add to the C compiler
 - OpenMP_CXX_FLAGS - flags to add to the CXX compiler
 - OPENMP_FOUND - true if openmp is detected
- Old version do not have Fortran flags!



CMake example for C

```
...  
  
if(OpenMP_Build)  
  find_package(OpenMP)  
  
  if (OPENMP_FOUND)  
    set (CMAKE_C_FLAGS "${CMAKE_C_FLAGS} ${OpenMP_C_FLAGS}")  
  endif (OPENMP_FOUND)  
endif(OpenMP_Build)  
  
...
```



CMake example for Fortran

```
...  
  
if(OpenMP_Build)  
  find_package(OpenMP)  
  
  if (OPENMP_FOUND)  
    set (CMAKE_Fortran_FLAGS "${CMAKE_Fortran_FLAGS} ${OpenMP_Fortran_FLAGS}")  
  endif (OPENMP_FOUND)  
endif(OpenMP_Build)  
  
...
```



Running OpenMP programs

- Make sure the compiler module used for building the executable is loaded (shared libraries!)
- Parallel jobs need to run inside the batch submission system
- Batch submission systems are vital to get consistent runtimes
- Expect a similar set-up on any well managed service

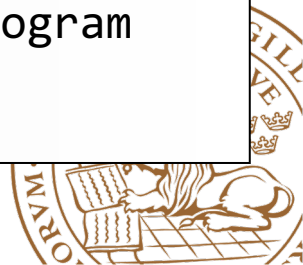


Standard OpenMP job

Run executable: processor_omp

```
#!/bin/bash
#SBATCH -c 5                # number of cpus per task
#SBATCH -t 00:05:00        # job-time - here 5 min
#SBATCH -J data_process    # name of job
#SBATCH -o process_omp_%j.out # output file
#SBATCH -e process_omp_%j.err # error messages

cat $0
ml purge > /dev/null 2>&1
module load <compiler/version> # replace as needed
export OMP_NUM_THREADS=$SLURM_CPUS_PER_TASK # set nr. threads
./processor_omp                 # run the program
```



OpenMP job in course reservation

Run executable: processor_omp

```
#!/bin/bash
#SBATCH -c 5                                # number cpus per task
#SBATCH -t 00:05:00                        # job-time - here 5 min
#SBATCH -J data_process                    # name of job
#SBATCH -A SNIC2022-22-140                # course project
#SBATCH --reservation=openmp-course-dayX  # access reserved nodes
#SBATCH -o process_omp_%j.out              # output file
#SBATCH -e process_omp_%j.err              # error messages

cat $0
ml purge > /dev/null 2>&1
module load <compiler/version>             # replace as needed
export OMP_NUM_THREADS=$SLURM_CPUS_PER_TASK # set nr. threads
./processor_omp                             # run the program
```



Running different number of threads than cores specified

- Can be useful for benchmarking
 - No disturbance from other jobs
 - **Remark:** You pay more cores than you use

- E.g.: to specify to run four threads specify

```
export OMP_NUM_THREADS=4
```

- Specify before starting your executable
- Make sure you asked for **more processors** than specified



Submission, queue monitoring and modification

INTERACTING WITH SLURM



Submission with sbatch

- Use sbatch to submit your job script to the job-queue
- Example:

```
[fred@alarik Timetest]$ sbatch runjob.sh  
Submitted batch job 7197
```

- Submit script “runjob.sh”
- Successful submission returns a job-id number (7197)



Monitoring the queue with squeue

- Use **squeue** to monitor the job queue

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
7303	snic	hybrid_n	fred	PD	0:00	32	(Priority)
7302	snic	hybrid_n	fred	PD	0:00	32	(Priority)
7301	snic	hybrid_n	fred	PD	0:00	32	(Resources)
7304	snic	preproce	karl	PD	0:00	6	(Priority)
7300	snic	hybrid_n	fred	R	0:24	32	an[001-032]
7305	snic	preproce	karl	R	0:37	6	an[081-086]
7306	snic	hybrid_n	fred	R	0:37	6	an[081-086]
7307	snic	testsimu	sven	R	0:07	1	an081

- Typically lots of output – use options of squeue to filter



Options of squeue

- Showing jobs for a specific user

```
squeue -u fred
```

will show the jobs of user “fred” only

- Option --start gives the estimated job start time
 - This can shift in either direction



Deleting jobs with scancel

- You can cancel a queued or running job
- Determine job-id, e.g. with squeue
- Use scancel

```
scancel 7103
```

- terminates job 7103, if running
- removes from the queue



Summary

- Kebnekaise's HPC2N hardware
- Building an OpenMP executable on Kebnekaise
- Running the executable using SLURM

