



LUND
UNIVERSITY



HPC2N, UmU

Data for parallel regions

Joachim Hein

LUNARC & Centre of Mathematical Sciences
Lund University

Pedro Ojeda May
HPC2N
Umeå University



Overview

- Data
 - What is private data
 - What is shared data
 - How to control which is which
- Race conditions
 - Basic constructs to avoid data races



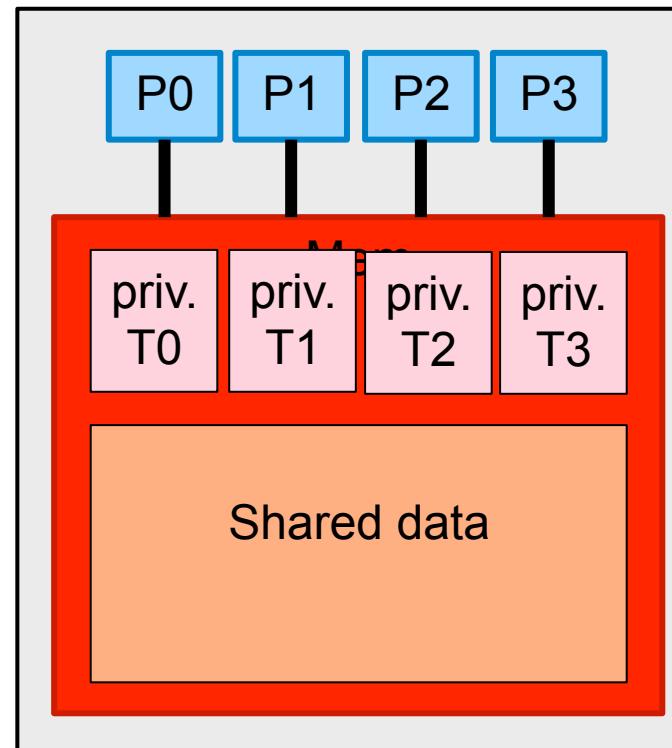
SHARED & PRIVATE DATA

Lund University / Faculty / Department / Unit / Document / Date



Private and shared data

- In parallel region
 - Shared and private data
- Shared data
 - Every thread can access (potential for conflict)
 - Unchanged on entry to par. region
 - Survives after end of par. Region
- Private data
 - Each thread: **own private copy**
 - Normally **uninitialised** at begin of parallel region
 - Contents typically **lost** when parallel region finishes



Fortran: Clauses to control data sharing

- For data declared before start of parallel region:
 - Use clause `shared` to declare a data structure as shared
 - Use clause `private` to declare a data structure as private

```
integer :: a=5, b  
  
!$omp parallel &  
 !$omp shared(a) private(b)  
 b = a &  
 + omp_get_thread_num()  
 print *, result=",b  
 !$omp end parallel
```



C: Clauses to control data sharing

- For data declared before start of parallel region:
 - Use clause `shared` to declare a data structure as shared
 - Use clause `private` to declare a data structure as private

```
int a, b;  
a = 5;  
#pragma omp parallel \  
shared(a) private(b)  
{  
    b = a +  
        omp_get_thread_num();  
    printf("%i\n", b);  
}
```



Private data

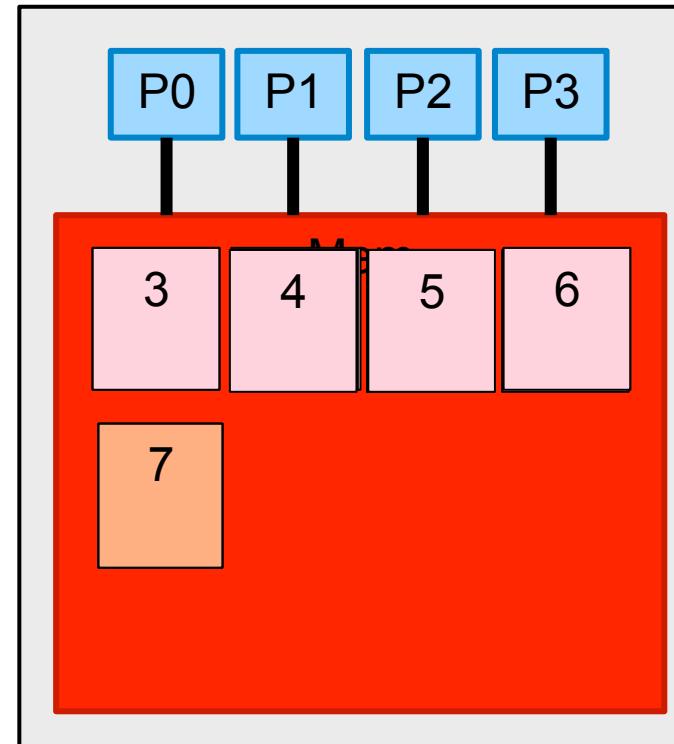
- Typically control variables, examples include
 - Thread identification
 - Loop indices
 - Variables internal to the algorithm
- Most variables declared inside a parallel region are private
 - Declared inside the block (C/C++)
 - In subroutine/function called from inside parallel region
 - **Exceptions include:**
 - static (C/C++) or save (Fortran) variables
 - file scope (C/C++) or COMMON blocks
 - Variables passed by reference inherit
- In Fortran: **Care needed** with COMMON and EQUIVALENCE



Example (Fortran): Memory movements for private data

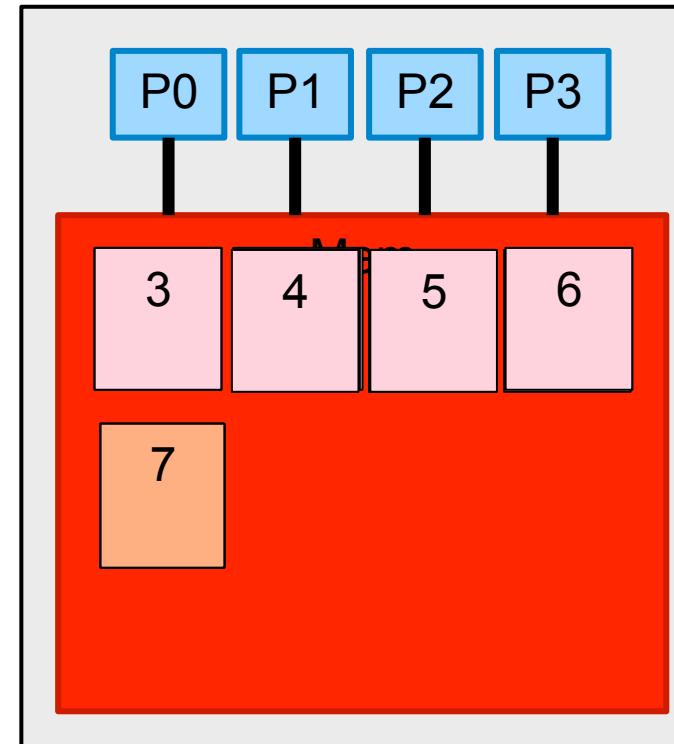
```
integer b  
b=5  
  
!$OMP parallel &  
!$OMP      private(b)  
b=get_omp_thread_num()  
b= b+3  
!$OMP end parallel
```

b=7



Example (C-code): Memory movements for private data

```
int b;  
b=5;  
#pragma omp parallel \  
    private(b)  
{  
    b=get_omp_thread_num();  
    b+=3;  
}  
b=7;
```



Shared Data

- Majority of the data
 - Typically: large data structures (e.g. arrays)
- Keeps its value on
 - Entry to parallel region
 - Exit from parallel region
- Every thread can access (read and/or write) the data
 - Save: when read by multiple threads (only read)
 - **Danger looms when:**
 - Multiple threads access the same memory location
 - At least one of these is a write access
 - Easily results in a “*Racecondition*”



Fortran-example: Initialisation of a Vector

$$v_i = 4i$$

```
integer, parameter:: vleng=120
integer:: vect(vleng), myNum, start, fin, i

 !$omp parallel shared(vect) &
 !$omp private(myNum, start, fin, i)
     myNum = vleng/omp_get_num_threads()
     start = 1 + omp_get_thread_num()      * myNum
     fin   =      (omp_get_thread_num() + 1) * myNum
     do i = start, fin
         vect(i) = 4*i ! threads write diff. elements
     enddo
 !$omp end parallel
```



C-example: Initialisation of a Vector

$$v_i = 4i$$

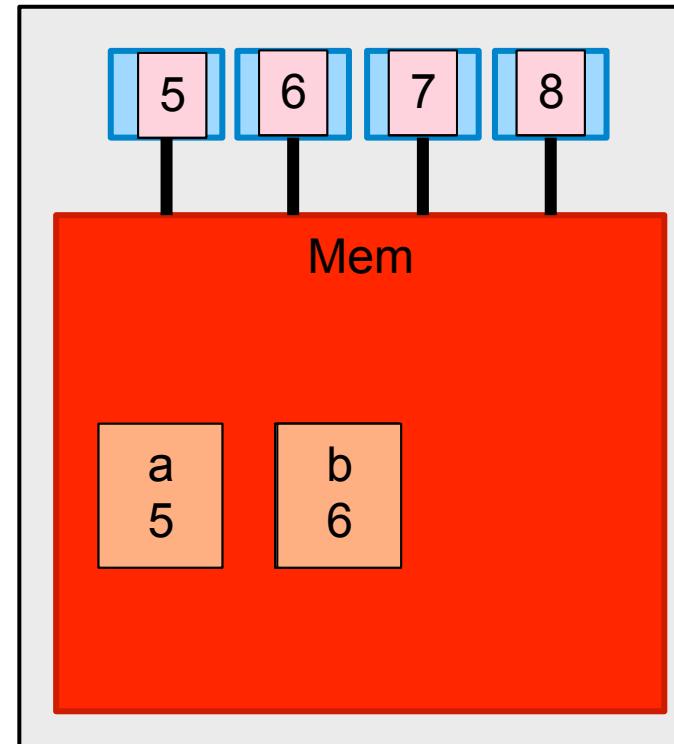
```
const int vleng=120;  
int vect[vleng], myNum, start, fin, i;  
  
#pragma omp parallel shared(vect) \  
private(myNum, start, fin, i)  
{  
    myNum = vleng/omp_get_num_threads();  
    start = omp_get_thread_num() * myNum;  
    fin = start + myNum;  
    for (i = start; i < fin; i++)  
        vect[i] = 4*i; // tasks write diff. elements  
}
```



Fortran example: Write conflict for shared data

```
integer :: a, b  
  
a=5  
  
!$OMP parallel &  
!$OMP    shared(a,b)  
    b = a +  
        omp_get_thread_num()  
    print *, "updated"  
    print *, "my b:", b  
!$OMP end parallel
```

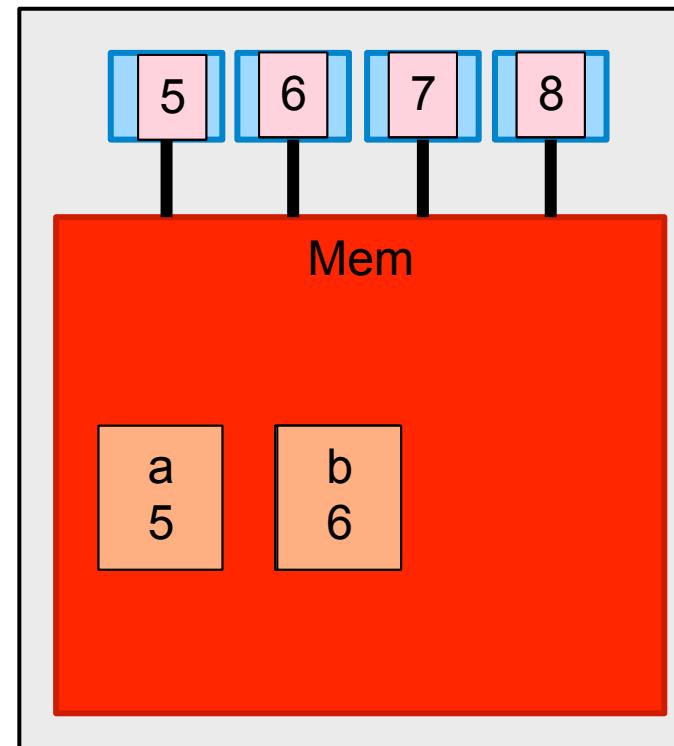
- final b-value: “random”
- Individual tasks might print b before b has final value!



C-example: Write conflict for shared data

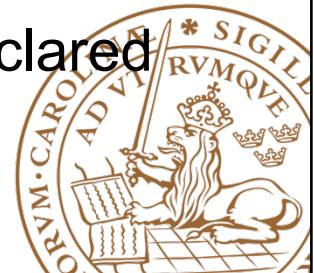
```
int a, b;  
  
a=5;  
  
#pragma omp parallel \  
shared(a,b)  
{  
    b = a +  
        omp_get_thread_num();  
    printf("updated\n");  
    printf("my b: %i", b);  
}
```

- final b-value: “random”
- Individual tasks might print b before b has final value!



Default clause

- Can be used on a `parallel` or `task` construct
- Determines data sharing of implicitly determined variables
 - In C
`default(shared|none)`
 - In Fortran
`default(shared|none|private|firstprivate)`
- For `parallel` construct, if no `default` clause supplied,
`default(shared)` applies
- **Recommendation:** `default(none)` typically good idea!
 - All variables accessed in parallel region must be declared as `shared`, `private`, ...



barrier, critical & atomic

FIXING DATA RACES



Data races

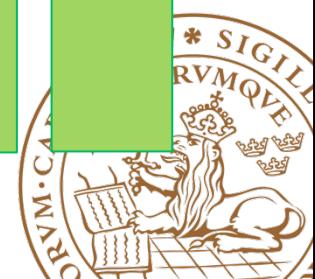
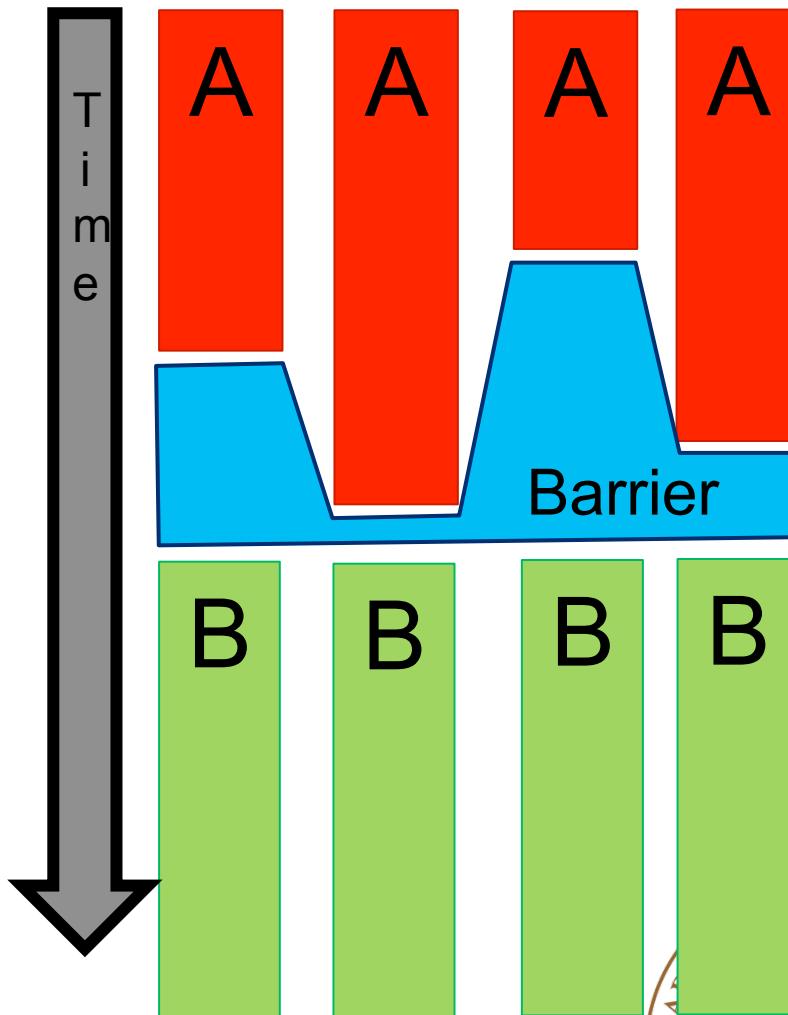
- Data races occur easily when:
 - Data is (to be) **written** to **shared memory location** on one thread
 - Accessed (read and/or write) by **another** thread
- Problems include:
 - Has writing thread reached the write statement?
 - If so, is the data in the memory system (cache, main memory) or still held in a processor register?
- Constructs to deal with data races impact code performance
 - Though we have no interest in “fast garbage”!!!



Barrier and synchronisation

- Barrier enforces synchronisation

```
#pragma omp barrier  
!$omp barrier
```
- All threads wait for the last one to arrive
- Registers are flushed to memory system
- All threads must have the barrier in their line of execution
 - Otherwise: **dead-lock!**



Example: Data race in matrix transpose

```
!$omp parallel default(none) &
 !$omp private(mysize,tid,i,j) shared(matrix,mtrans)
   tid = omp_get_thread_num()
   mysize = nszie/omp_get_num_threads()
   do j = 1 + tid*mysize, (tid+1)*mysize
     do i = 1, nszie
       matrix(i,j) = 1000.0 * j + i
     enddo
   enddo
 !$omp barrier
   do j = 1 + tid*mysize, (tid+1)*mysize
     do i = 1, nszie
       mtrans(i,j) = matrix(j,i)
     enddo
   enddo
 !$omp end parallel
```



Example: Data race in matrix transpose

```
#pragma omp parallel default(none) \
    private(mysize, tid, i, j) shared(matrix, mtrans)
{ tid = omp_get_thread_num();
  mysize = nsize/omp_get_num_threads();

  for (i = tid*mysize; i < (tid+1)*mysize; i++)
    for (j = 0; j < nsize; j++)
      matrix[i][j] = 1000.0 * j + i;

  #pragma omp barrier
  for (i = tid*mysize; i < (tid+1)*mysize; i++)
    for (j = 0; j < nsize; j++)
      mtrans[i][j] = matrix[j][i];
}
```



Critical regions

- Critical region: protect update of share memory location
- Only one thread executes critical region at the time

```
#pragma omp critical (name)
{
    code-block
}
```

- Name is optional
 - Single thread in all regions of same name
 - Single thread in all unnamed regions
- Implies a register flush at entrance/exit
- Also useful to execute non-thread-safe functions
- Performance penalty due to serialisation ☹



Critical regions

- Critical region: protect update of share memory location
- Only one thread executes critical region at the time

```
!$omp critical (name)
code-block
!$omp end critical (name)
```

- Name is optional
 - Single thread in all regions of same name
 - Single thread in all unnamed regions
- Implies a register flush at entrance/exit
- Also useful to execute non-thread-safe functions
- Performance penalty due to serialisation ☹



Fortran syntax for critical region

- In Fortran a start and end mark is required

```
!$omp critical (name)
    code-block
!$omp end critical (name)
```

- Behaviour as in C, see above



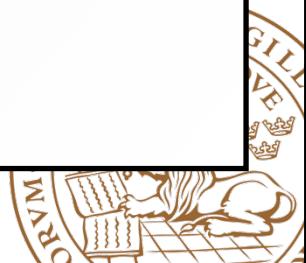
$$\sum_{k=0}^{n-1} e^k$$

Example for use of **critical**

```
sum=0.0_dpr
!$omp parallel default(none) &
!$omp    shared(sum) private(tid, cont)

    tid = omp_get_thread_num()
    cont = func( tid )
!$omp critical (exp_up)
    sum = sum + cont
    print *, tid,: c=", cont, " s=", sum
!$omp end critical (exp_up)

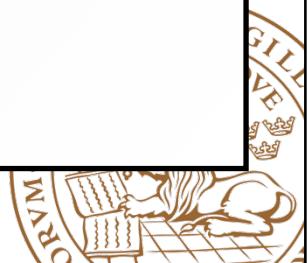
!$omp end parallel
```



$$\sum_{k=0}^{n-1} e^k$$

Example for use of **critical**

```
sum=0.0;  
#pragma omp parallel default(none) \  
    shared(sum) private(tid, cont)  
{  
    tid = omp_get_thread_num();  
    cont = func( tid );  
    #pragma omp critical (exp_up)  
    {  
        sum += cont;  
        printf("%i: c=%f s=%f\n", tid, cont, sum);  
    }  
}
```



Atomic

- Light-weight alternative to `critical` in simple cases
- Works with simple statements only
- Can use special hardware instructions if exist
- Flushes “protected” variable on entry and exit
- Four different versions, from OpenMP 3.1
 - `read`
 - `write`
 - `update`
 - `capture`
- OpenMP 4.0: adding `seq_cst` to `atomic` flushes all variables
 - Important for controlling instruction reordering
 - Example: implementing a lock



Atomic read of shared variable

- Protects only the reading of scalar intrinsic variable `x`
- Flushes `x` on entry and exit
- Fortran:

```
!$omp atomic read  
v = x
```

- C:

```
#pragma omp atomic read  
v = x;
```



Atomic write of shared variable

- Protects only the writing of scalar intrinsic variable `x`
- No protection of evaluation of `expr` on the right hand side
- Flushes `x` on entry and exit
- Example expressions:
 - `x = 5;`
 - `x = v;`
 - `x = func(a);`
- Fortran:

```
!$omp atomic write
x = expr
```
- C:

```
#pragma omp atomic write
x = expr;
```



Atomic update of shared variable

- Only protects the update of `x`, not the right hand call to `func()`

```
x += func(a);  
x = x + func(a)
```
- Works with simple statements only
- Can use special hardware instructions if exist
- Flushes update variable on entry and exit
- **Rem:** Only atomic operation prior to OpenMP 3.1



Fortran statements: atomic update

- Examples:

```
!$omp atomic update  
  x = x + 1  
  
!$omp atomic update  
  x = x + f(a)
```

- Allowed operations:
 - $x = x \operatorname{operator} expr$
 - $x = expr \operatorname{operator} x$
 - $x = \operatorname{intr_proc}(x, \operatorname{expr_list})$
 - $x = \operatorname{intr_proc}(\operatorname{expr_list}, x)$
 - x scalar, intrinsic type
 - $\operatorname{operator}$ is one of:
 - $+, *, -, /,$
 - $.AND., .OR., .EQV., .NEQV.$
 - $\operatorname{intr_proc}$ is one of
 - $\operatorname{MAX}, \operatorname{MIN}, \operatorname{IAND}, \operatorname{IOR}, \operatorname{IEOR}$
- The evaluation of $f(a)$ is **not** protected – use `critical` if needed!
 - The update is optional – consistency with old OpenMP standard



F90-example: Vector norm

$$\sqrt{\sum_i v(i) * v(i)}$$

```
norm = 0.0D0
 !$omp parallel default(none) &
 !$omp shared(vect,norm) private( myNum, i, lNorm)
 lNorm = 0.0D0
 myNum = vlen/omp_get_num_threads() !local size
 do i = 1 + myNum * omp_get_thread_num(), &
       myNum *(1+omp_get_thread_num())
   lNorm = lNorm + vect(i)*vect(i)
 enddo
 !$omp atomic update
 norm = norm + lNorm
 !$omp end parallel
 norm = sqrt(norm)
```



C statements: atomic update

Examples:

```
#pragma omp atomic update  
    x++;  
  
#pragma omp atomic update  
    x += f(a);
```

- Evaluation of $f(a)$ is **not** protected – use `critical` if needed!
- The update is optional – consistency with old OpenMP standard

- Allowed operations:

- $x \ binop= \ expr;$
- $x++;$
- $++x;$
- $x--;$
- $--x;$
- $x = x \ binop \ expr;$

- x lvalue, scalar

- $binop$ is one of:

- $+, *, -, /, \&, ^, |, <<, >>$



C-example: Vector norm

$$\sqrt{\sum_i v(i) * v(i)}$$

```
norm = 0.0;  
#pragma omp parallel default(none) \  
    shared(vect,norm) private( myNum, i, lNorm)  
{ lNorm = 0.0;  
    myNum = vlen/omp_get_num_threads(); //local size  
    for (i = myNum*omp_get_thread_num();  
         i < myNum*(1+omp_get_thread_num()); i++)  
        lNorm += vect[i]*vect[i];  
#pragma omp atomic update  
    norm += lNorm;  
}  
                                // synchronise at end parallel  
norm = sqrt(norm);
```



Atomic capture of a shared variable

- Atomic capture:
 - Update a shared variable
 - Keep a thread private copy of either (but not both)
 - Old value
 - New value
- Restrictions apply



C statements: atomic capture

```
#pragma omp atomic capture  
statement or strct. blk.
```

Allowed statements (OpenMP 4.0):

v = x++;

v = x--;

v = ++x;

v = --x;

v = x binop= expr;

v = x = x binop expr;

v = x = expr binop x;

Allowed structured blocks:

{v = x; x binop= expr;}

{x binop= expr; v = x;}

{v = x; x = x binop expr;}

{v = x; x = expr binop x;}

{x = x binop expr; v = x;}

{x = expr binop x; v = x;}

{v = x; x = expr;}

{v = x; x++;}

{v = x; ++x;}

{++x; v = x;}

{x++; v = x;}

{v = x; x--;}

{v = x; --x;}

{--x; v = x;}

{x--; v = x;}



Fortran statements: atomic capture

```
!$omp atomic capture  
    update-statement  
    capture statement  
 !$omp end atomic
```

or

```
!$omp atomic capture  
    capture statement  
    update-statement  
 !$omp end atomic
```

- Allowed update statements:
 $x = x \operatorname{operator} expr$
 $x = expr \operatorname{operator} x$
 $x = \operatorname{intr_proc}(x, \operatorname{expr_list})$
 $x = \operatorname{intr_proc}(\operatorname{expr_list}, x)$

- Allowed capture statements:
 $v = x$



Summary

- Private data
- Shared data
- Prevent race condition
 - `barrier`
 - `critical`
 - `atomic`
- Examples also showed parallelisation strategies

