

Worksharing and scheduling

Pedro Ojeda & Joachim Hein

High Performance Computing Center North &
Lund University

Overview

- **Worksharing constructs** allow easy distribution of the work onto the threads
 - Loop construct
 - Easy distributing of loops onto threads
 - Avoiding load imbalance in distributed loop using the `schedule` clause
 - `workshare` construct
 - Parallelisation of Fortran array syntax
 - `sections` construct
 - Distributing independent code blocks
 - Modern alternative: `task`



DISTRIBUTING LOOPS



Introduction for the loop construct

- Distributing large loops: typical target for OpenMP
- As seen in previous lecture: Can be accomplished with the features discussed
- Requires some management code:
 - Number of iterations per thread (may be unequal)
 - Starting index of current thread
 - Final index of current thread
- OpenMP offers “*loop construct*” to ease loop parallelisation
 - Convenience
 - Maintainability



The loop construct

- Distributes the following loop (C/C++/Fortran) onto threads
- The iteration variable automatically `private`
- Determines for you (without management code):
 - Number of iterations per thread
 - Start index of current thread
 - Final index of current thread
- Registers are flushed to memory at exit, unless `nowait`
 - **No** `flush` on entry!
- Offers mechanisms to balance the load for a number of situations



Loop construct in Fortran

- Works on Fortran standard compliant `do-construct`
 - not: `do while`
 - not: `do without loop control`
- `$omp end do` not required, optional

```
!$omp parallel &  
!$  shared(...) &  
!$  private(...)  
  
    !$omp do  
    do i=1, N  
        loop-body  
    end do  
  
!$omp end parallel
```



F90-example: Vector norm

Managing loop yourself

$$\sqrt{\sum_i v(i) * v(i)}$$

```
norm = 0.0D0
!$omp parallel default(none) &
!$omp    shared(vect,norm) private( myNum, i, lNorm)
    lNorm = 0.0D0
    myNum = vleng/omp_get_num_threads() !local size
    do i = 1 + myNum * omp_get_thread_num(), &
        myNum * (1+omp_get_thread_num())
        lNorm = lNorm + vect(i)*vect(i)
    enddo
!$omp atomic update
    norm = norm + lNorm
!$omp end parallel
norm = sqrt(norm)
```



Example: Vector norm

Loop construct

$$\sqrt{\sum_i v(i) * v(i)}$$

```
norm = 0.0d0;
!$omp parallel default(none) &
!$omp    shared(vect,norm) private( i, lNorm )
    lNorm = 0.0d0

!$omp do
    do i = 1, vleng                //same as serial case
        lNorm = lNorm + vect(i) * vect(i)
    enddo

!$omp atomic update
    norm = norm + lNorm
!$omp end parallel
    norm = sqrt(norm)
```



Loop construct in C

- Limited to “canonical” loops:
- First argument assignment to
 - int
 - pointer
 - random-access-iterator-type (C++)
- Second arg. comparison:
 - using `<=`, `<`, `>`, `>=`
- Third arg: increment
 - `i++`, `++i`, `i--`, `--i`
 - `i+=inc`, `i-=inc`
 - `i=i+inc`, `i=inc+i`, `i=i-inc`
- All bounds, increments: **loop-invariant**

```
#pragma omp parallel \  
    shared(...) \  
    private(...)  
  
#pragma omp for  
for (i=0; i<N; i++)  
{  
    loop-body  
}
```



Example: Vector norm

Managing loop yourself

$$\sqrt{\sum_i v(i) * v(i)}$$

```
norm = 0.0;
#pragma omp parallel default(none) \
    shared(vect,norm) private( myNum, i, lNorm)
{ lNorm = 0.0;
  myNum = vlen/omp_get_num_threads(); //local size
  for (i = myNum * omp_get_thread_num();
       i < myNum * (1+omp_get_thread_num()); i++)
    lNorm += vect[i]*vect[i];
  #pragma omp atomic update
  norm += lNorm;
}
norm = sqrt(norm);
```



Example: Vector norm

Loop construct

$$\sqrt{\sum_i v(i) * v(i)}$$

```
norm = 0.0;
#pragma omp parallel default(none) \
    shared(vect,norm) private( i, lNorm )
{ lNorm = 0.0;
  #pragma omp for
  for (i = 0; i < vleng; i++) //same as serial case

      lNorm += vect[i]*vect[i];
  #pragma omp atomic update
  norm += lNorm;
}
norm = sqrt(norm);
```



Parallel loop construct in Fortran

- Shorthand if parallel region being a loop construct

```
!$omp parallel do
  do i, N
    loop-body
  enddo  // parallel region ends here!
```

- No `!$omp end parallel do` required (optional)
- Features of parallel region and normal loop construct apply similarly



Parallel loop construct in C

- Shorthand if parallel region being a loop construct

```
#pragma omp parallel for
for (int i; i<N; i++)
{
    loop-body
} // parallel reg. & loop constr. end here!
```

- Features of parallel region and normal loop construct apply similarly



Loop reordering and data dependency

- In parallel loop: Iterations executed in *different order* from serial code
- Correct result only if current iteration *independent* of previous iteration (data dependency)
- If data dependency
 - modify/change algorithm
 - serialise relevant part of the loop using special OpenMP features (later in course)
 - execute loop serial

- Example for dependency:

```
a[0]=0;  
for (i=1; i<N; i++)  
    a[i] = a[i-1] + i;
```

- Possible fix (algor. change):

```
for (i=0; i<N; i++)  
    a[i]=0.5*i*(i+1);
```



SCHEDULING LOOP ITERATIONS



Work per loop iteration

- So far assumed: Same work for each loop iterations
- Not always the case, e.g.:

- Summing over triangular area

```
for (i=0; i<N; i++)
```

```
    for (j=0; j<i+1; j++)
```

- Loop body iterates until required accuracy achieve
- Often cause for *load-imbalance*:
 - Some threads finished while others still work
 - ➔ Poor performance!!!
- Dealing with such problems is typically easier in shared memory than in distributed memory programming



Schedule clause

- To help load balance in a loop construct:

```
schedule(kind, [chunk_size])
```

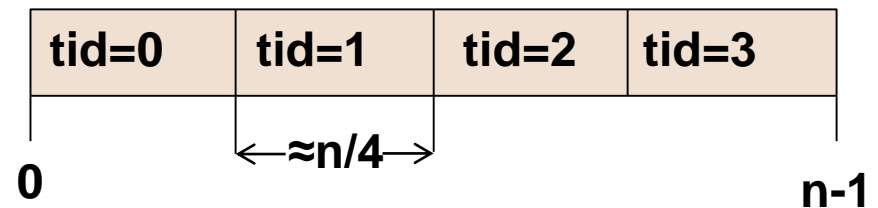
- Default schedule implementation dependent (OpenMP 3.0)
- Choices for `kind`:
 - `static`
 - `dynamic`
 - `guided`
 - `auto`
 - `runtime`



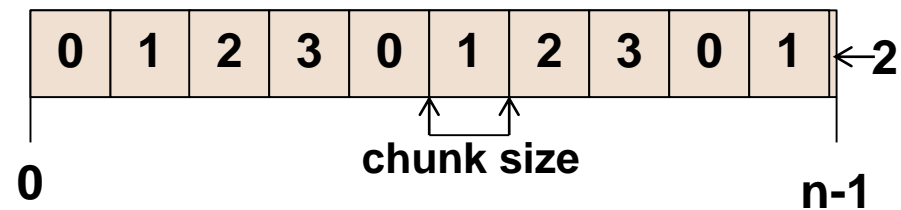
Static scheduling

- Divide iteration count into chunks of equal size
 - Last chunk smaller if needed
- Thread assignment:
“Round robin”
- Default chunk size: divide iteration count by number of threads
- Least overhead compared to other schedules

Default static schedule

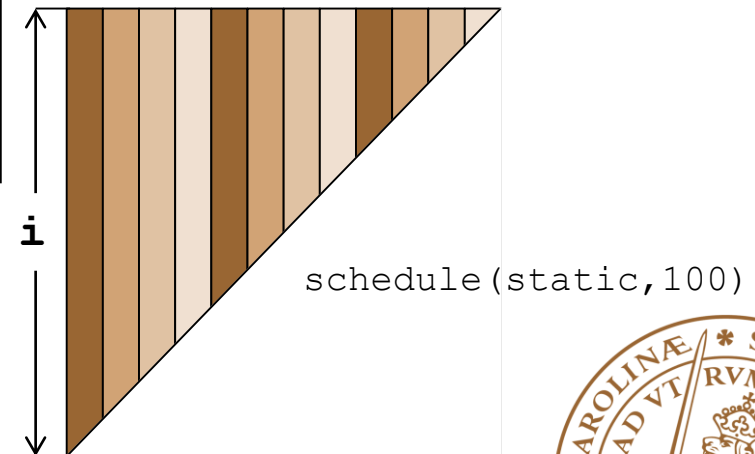
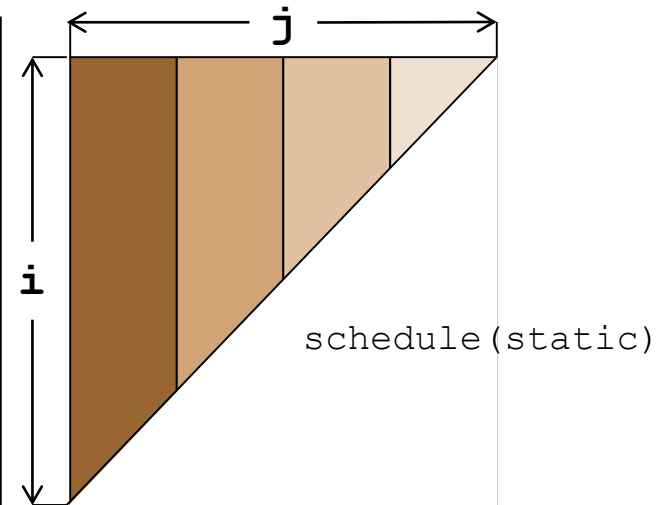


Static schedule



Example: Summation over triangular area

```
!$omp parallel do &  
!$omp private(i,j) shared(a) &  
!$omp schedule(static, 100)  
  do j=1, 1200  
    do i=j+1, 1200  
      a(i,j) = func(i,j)  
      a(j,i) = -a(i,j)  
    enddo  
  enddo
```



- default: max: 7/16 area
- static,100: max: 5/16 area
- smaller chunks: better balance
- more chunks: larger overhead



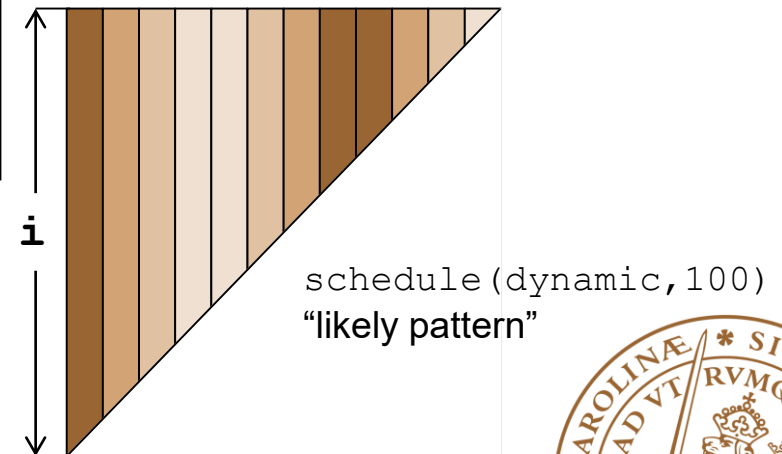
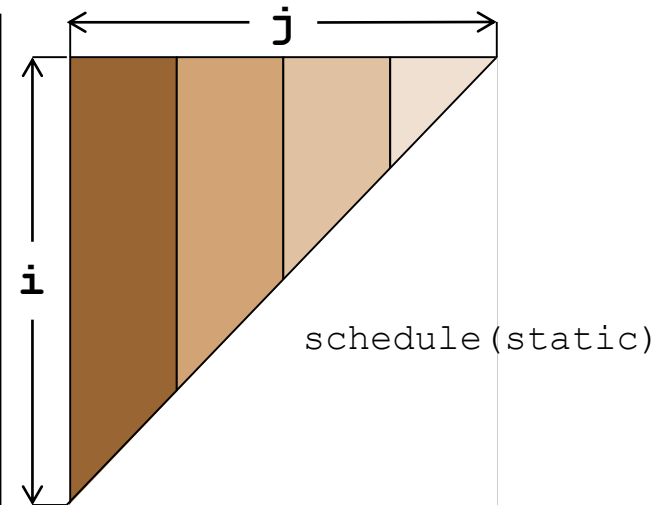
Dynamic scheduling

- Loop split into work packages of `chunk_size` iterations
- Each thread requests new work package once done with present
- Default `chunk_size`: 1 iteration



Example: Summation over triangular area

```
!$omp parallel do &  
!$omp private(i,j) shared(a) &  
!$omp schedule(dynamic, 100)  
  do j=1, 1200  
    do i=j+1, 1200  
      a(i,j) = func(i,j)  
      a(j,i) = -a(i,j)  
    enddo  
  enddo
```



- default: max: 7/16 area
- dynamic, 100: max: ≈ 0.27 area
- better balance than static
- larger overhead than static



Guided scheduling

- Similar to `dynamic`: Task request new work package once done
- Work package size proportional to
number of unassigned iterations / number of threads
- But never smaller than `chunk_size` (unless last WP)
- Default `chunk_size = 1`
- Idea: Prevent expensive work packages at the end



Schedules auto and runtime

- For `auto` the implementation decides
- For `runtime` the schedule can be controlled:
 - set schedule using functions `omp_set_schedule()`
 - this is OpenMP 3.0
 - Example: `omp_set_schedule(omp_sched_static, 10)`
 - read environment variable `OMP_SCHEDULE`
 - `export OMP_SCHEDULE="guided,4"`
 - `setenv OMP_SCHEDULE "guided,4"`
- Do not specify `chunk_size` with `auto` or `runtime`



MULTIPLE LOOP PARALLELISATION



Simple example with two nested loops

- 3 basic options to parallelise
- Depends which one is best

```
do j=1, 3

    do i=1, 4

        a(i,j) =
            expensiveFunc(i,j)
    enddo
enddo
```



Simple example with two nested loops

- Distribute the j -loop
- Maximally 3 work-packages

```
!$omp parallel do
do j=1, 3

    do i=1, 4

        a(i,j) =
            expensiveFunc(i,j)
    enddo
enddo
```



Simple example with two nested loops

- Distribute i -loop
- Now four work packages
- Parallel before j -loop
 - better performance
 - needs private i
- Starts loop construct 3 times
- (more) cache line conflict when writing to a

```
!$omp parallel private(j)
do j=1, 3

!$omp do
do i=1, 4

    a(i,j) =
        expensiveFunc(i,j)
enddo

!$omp end do
enddo

!$omp end parallel
```



Simple example with two nested loops

- Use collapse clause
 - specify number of loops to collapse
 - **OpenMP 3.0**
- Distribute both loop
 - Creates single loop
 - Schedule as specified
 - default in this case
- Now: 12 work-packages
- (more) cache line conflict when writing to a

```
!$omp parallel
!$omp do collapse(2)
do j=1, 3
  do i=1, 4

    a(i,j) =
      expensiveFunc(i,j)

  enddo
enddo
```



Simple example with two nested loops

- 3 basic options to parallelise
- Depends which one is best

```
for (int i=0; i<3; i++)  
{  
  
    for (int j=0; j<4; j++)  
    {  
        a[i][j] =  
            expensiveFunc(i,j);  
    }  
}
```



Simple example with two nested loops

- Distribute the i -loop
- Maximally 3 work-packages

```
#pragma omp parallel for
for (int i=0; i<3; i++)
{
    for (int j=0; j<4; j++)
    {
        a[i][j] =
            expensiveFunc(i,j);
    }
}
```



Simple example with two nested loops

- Distribute j -loop
- Now four work packages
- Parallel before i -loop
 - better performance
 - needs private i
- Starts loop construct 3 times
- (more) cache line conflict when writing to a

```
#pragma omp parallel
{
  for (int i=0; i<3; i++)
  {
    #pragma omp for
    for (int j=0; j<4; j++)
    {
      a[i][j] =
        expensiveFunc(i,j);
    }
  }
}
```



Simple example with two nested loops

- Use collapse clause
 - specify number of loops to collapse
 - **OpenMP 3.0**
- Distribute both loop
 - Creates single loop
 - Schedule as specified
 - default in this case
- Now: 12 work-packages
- (more) cache line conflict when writing to a

```
#pragma omp parallel
#pragma omp for collapse(2)
for (int i=0; i<3; i++)
{
    for (int j=0; j<4; j++)
    {
        a[i][j] =
            expensiveFunc(i,j);
    }
}
```



WORKSHARE IN FORTRAN



Workshare construct for Fortran

- For Fortran OpenMP provides `workshare`

- This allows distribution of

- Fortran array syntax

`a(1:n, 1:m) = b(1:n, 1:m) + c(1:n, 1:m)`

- Fortran: `FORALL`, `WHERE`



Example for workshare

```
!$OMP PARALLEL SHARED(n, a, b, c)
!$OMP WORKSHARE
  b(1:n) = b(1:n) + 1
  c(1:n) = c(1:n) + 2
  a(1:n) = b(1:n) + c(1:n)
!OMP END WORKSHARE
!OMP END PARALLEL
```

- OpenMP ensures there is no data race
 - b and c ready before assignment to a
- Can have user defined functions, if declared ELEMENTAL



Scalar assignment in workshare

```
      REAL AA (N,N) ,  BB (N,N) ,  CC (N,N) ,  DD (N,N)
      INTEGER SHR
!$OMP PARALLEL SHARED (SHR)
!$OMP WORKSHARE
      AA=BB
      SHR=1
      CC=DD *  SHR
!$OMP END WORKSHARE
!$OMP END PARALLEL
```

- Legal OpenMP
- Single thread performs scalar assign to SHR



Scalar assignment to **private** in workshare **ILLEGAL!!!**

```
      REAL AA (N,N) ,  BB (N,N) ,  CC (N,N) ,  DD (N,N)
      INTEGER PRI
!$OMP PARALLEL PRIVATE (PRI)
!$OMP WORKSHARE
      AA=BB
      PRI=1
      CC=DD * PRI
!$OMP END WORKSHARE
!$OMP END PARALLEL
```

- Single thread performs scalar assign to PRI
- Undefined on other threads 😞



SECTIONS

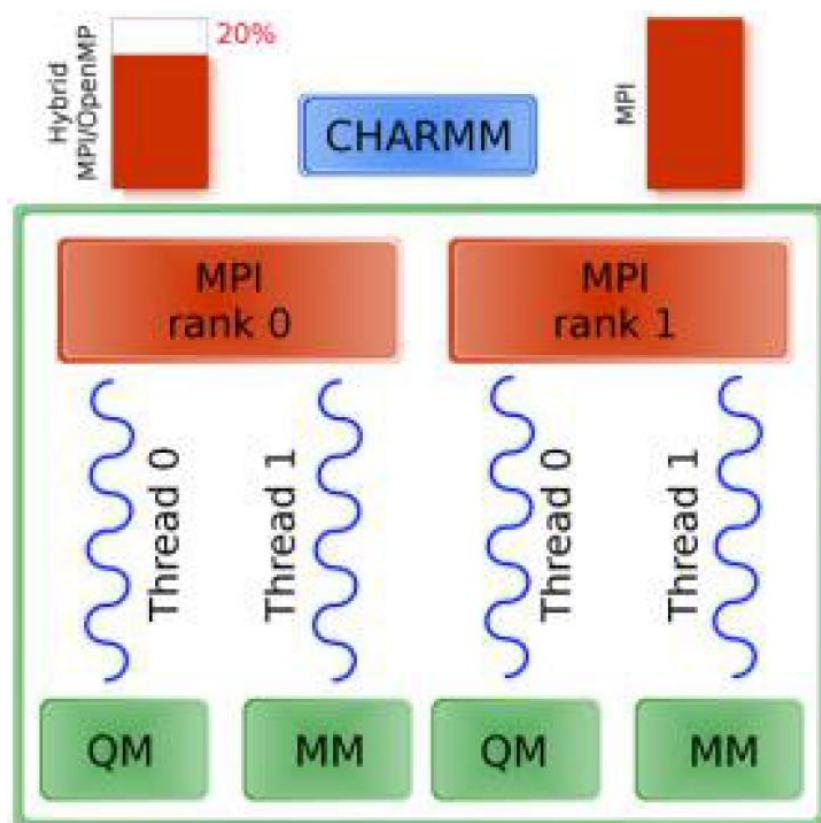


Sections construct

- Allows parallelisation when code blocks can be executed independently
 - Example: Initialisation of multiple data structures
- Easy way to get different tasks to execute different code
- Mismatch: Number of code blocks and number of threads
 - Individual threads might execute multiple code blocks
 - Not every thread gets a code block
- Danger of load imbalance
 - Code blocks have different amount of work
 - Mismatch block and thread number



Sections construct



Acceleration of Semiempirical QM/MM methods, JCTC, 13, 3525-3536 (2017)



Example for use of sections

```
#pragma omp parallel shared(a, b, N, M)
{
    #pragma omp sections
    {
        #pragma omp section
        { for(int i; i < N; i++)
            a[i] = i; }

        #pragma omp section
        { for(int i; i < M; i++)
            b[i] = initBmatrix(i,M); }
    }
}
```



Alternatively: parallel sections

```
#pragma omp parallel sections shared(a, b, N, M)
{
    #pragma omp section
    { for(int i; i < N; i++)
        a[i] = i; }
    #pragma omp section
    { for(int i; i < M; i++)
        b[i] = initBmatrix(i,M); }
}
```



Summary

- OpenMP loop construct
 - easy distribution of standard `do/for` loops
 - the `schedule` clause deals with many cases of load imbalance
- OpenMP `workshare` construct
 - distributes Fortran array syntax statements



Summary

- OpenMP loop construct
 - easy distribution of standard `do/for` loops
 - the `schedule` clause deals with many cases of load imbalance
- OpenMP `workshare` construct
 - distributes Fortran array syntax statements
- OpenMP `sections` construct
 - distributes independent code block on different threads

