

More on private data

Pedro Ojeda & Joachim Hein

High Performance Computing Center North &
Lund University


Outline

- Special versions of private data
 - firstprivate
 - lastprivate
 - reduction
 - **global storage:** threadprivate



Private and shared data

- In parallel region
 - Shared and private data
- Shared data
 - Unchanged on entry to par. region
 - Survives after end of par. Region
- Private data
 - Each thread: **own private copy**
 - Normally **uninitialised** at begin of parallel region
 - Contents typically lost when parallel region finishes
 - Though connection to values before/after often needed



Clause `firstprivate`

- Private variables are **not** initialised
- Clause `firstprivate`
 - declares variable(s) private
 - initialises each private copy with the value prior to the construct

Example local accumulation:

```
integer lsum=10
!$omp parallel &
!$omp  firstprivate(lsum)

    lsum = lsum &
        + omp_get_thread_num()
    print *, lsum
!$omp end parallel
```



Clause `firstprivate`

- Private variables are **not** initialised
- Clause `firstprivate`
 - declares variable(s) private
 - initialises each private copy with the value prior to the construct

Example local accumulation:

```
int lsum=10;
#pragma omp parallel \
    firstprivate(lsum)
{
    lsum +=
        omp_get_thread_num();
    printf("%i\n", lsum);
}
```



Fortran-example: Vector norm *private*

$$\sqrt{\sum_i v(i) * v(i)}$$

```
norm = 0.0
```

```
!$omp parallel default(none) \  
  shared(vect, norm) private(i, lNorm)  
    lNorm = 0.0  
    !$omp do  
      do i = 0, vleng  
        lNorm = lNorm + vect(i)**2  
      enddo  
      #pragma omp atomic update  
      norm += lNorm  
    !$omp end parallel  
    norm = sqrt(norm)
```



Fortran-example: Vector norm

firstprivate

$$\sqrt{\sum_i v(i) * v(i)}$$

```
norm = 0.0
lNorm = 0.0
!$omp parallel default(none) \
    shared(vect, norm) private(i) firstprivate(lNorm)

    !$omp do
    do i = 0, vleng
        lNorm = lNorm + vect(i)**2
    enddo
    #pragma omp atomic update
    norm += lNorm
!$omp end parallel
norm = sqrt(norm)
```



C-example: Vector norm *private*

$$\sqrt{\sum_i v(i) * v(i)}$$

```
norm = 0.0;
```

```
#pragma omp parallel default(none) \  
  shared(vect, norm) private(i, lNorm)  
{ lNorm = 0.0;  
  #pragma omp for  
  for (i = 0; i < vleng; i++)  
    lNorm += vect[i]*vect[i];  
  #pragma omp atomic update  
  norm += lNorm;  
}  
norm = sqrt(norm);
```



C-example: Vector norm

firstprivate

$$\sqrt{\sum_i v(i) * v(i)}$$

```
norm = 0.0;
lNorm = 0.0;
#pragma omp parallel default(none) \
    shared(vect, norm) private(i) firstprivate(lNorm)
{
    #pragma omp for
    for (i = 0; i < vleng; i++)
        lNorm += vect[i]*vect[i];
    #pragma omp atomic update
    norm += lNorm;
}
norm = sqrt(norm);
```




Clause lastprivate

- Clause `lastprivate`
 - Use with **loop**, **sections**
 - Variable private
 - In the end: assigns value from last iteration or section
 - Undefined if not set in last iteration/section
- Variables can be both: `firstprivate` & `lastprivate`

Example:

```
integer i, a
!$omp parallel do &
!$omp lastprivate(a)
do i=1, 100
    a=i+1
    func(a)
enddo
print *, "a=", a
! this prints: a=101
```




Clause `lastprivate`

- Clause `lastprivate`
 - Use with **loop**, **sections**
 - Variable private
 - In the end: assigns value from last iteration or section
 - Undefined if not set in last iteration/section
- Variables can be both: `firstprivate` & `lastprivate`

Example:

```
integer i, a
#pragma omp parallel \
    for lastprivate(a)
for (i=0; i<100; i++)
{ a=i+1;
  func(a);
}
printf("a=%i\n", a);
// this prints: a=100
```



Reduction variables

- Frequently needed: Reduction of private variables
 - E.g.: Averages of array values, scalar products
- We have done this before: example vector norm
 - used `atomic` to protect the update
- For a `reduction`, we have to specify
 - operation, e.g.: addition, multiplication, or, ...
 - one or more variables
 - A construct can have more than one `reduction`



Behavior of reduction

- The basic syntax

`reduction(operator : variable_list)`

- Variables specified in `reduction`:

- Private copy per thread

- Initialised with default matching on operator

- At the end of construct (e.g. parallel region)


- Value prior to construct combined with private copies
 - Using the specified operator for combining values
 - New value available after the construct



Example:

Memory movements for reduction

```
int b;  
b=5;  
#pragma omp parallel \  
    reduction(+:b)  
{  
    b+=omp_get_thread_num();  
}  
printf("%i\n", b);
```



Example:

Memory movements for reduction

```
integer b
```

```
b=5
```

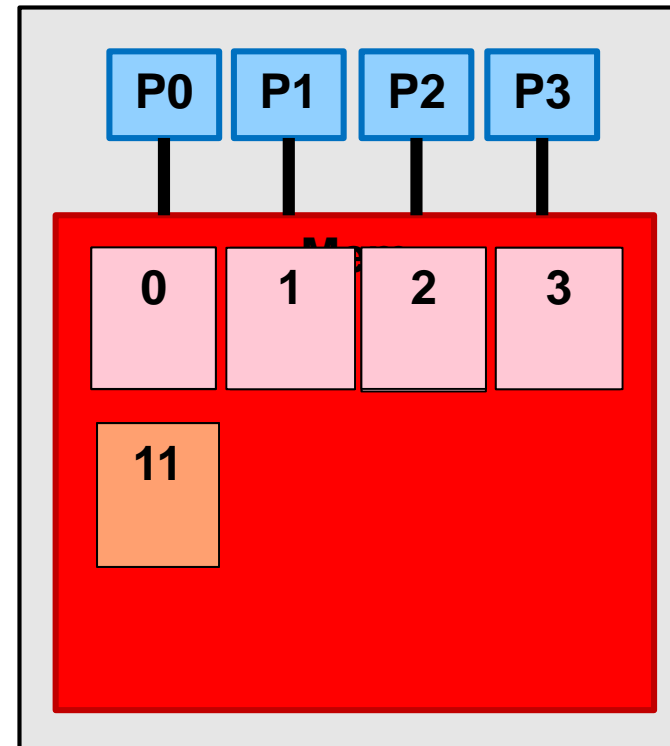
```
!$omp parallel &
```

```
!$omp  reduction(+:b)
```

```
  b=b+omp_get_thread_num()
```

```
!$omp end parallel
```


```
print *, b
```



Fortran example: Vector norm *atomic update*

$$\sqrt{\sum_i v(i) * v(i)}$$

```
norm = 0.0
lNorm = 0.0
!$omp parallel default(none) &
!$omp shared(vect,norm) private(i) firstprivate(lNorm)
!$omp do
    do i = 1, vleng
        lNorm = lNorm + vect(i)**2           ! priv. copy
    enddo
!$omp atomic update
    norm = norm + lNorm
!$omp end parallel
norm = sqrt(norm)                           ! comb. copies
                                              ! master copy
```



Fortran example: Vector norm *reduction*

$$\sqrt{\sum_i v(i) * v(i)}$$

```
norm = 0.0                                ! master copy
                                           ! lNorm gone

!$omp parallel default(none) &
!$omp shared( vect ) reduction( + : norm ) private(i)
!$omp do                                ! priv. copy=0
    do i = 1, vleng
        norm = norm + vect(i)**2        ! priv. copy
    enddo

!$omp end parallel                        ! comb. copies
norm = sqrt(norm)                        ! master copy
```



Fortran example: Vector norm *reduction, parallel do*

$$\sqrt{\sum_i v(i) * v(i)}$$

```
norm = 0.0
```

```
! master copy
```

```
!$omp parallel do default(none) &
```

```
!$omp shared( vect ) reduction( + : norm )
```

```
do i = 1, vleng
```

```
    norm = norm + vect(i)**2
```

```
! priv. copy
```

```
enddo
```

```
!$omp end parallel do
```

```
norm = sqrt(norm)
```

```
! master copy
```



Example: Vector norm *atomic update*

$$\sqrt{\sum_i v(i) * v(i)}$$

```
norm = 0.0;
lNorm = 0.0;
#pragma omp parallel default(none) \
    shared(vect, norm) private(i) firstprivate(lNorm)
{
    #pragma omp for
    for (i = 0; i < vleng; i++)
        lNorm += vect[i]*vect[i];
    #pragma omp atomic update
    norm += lNorm;
}
norm = sqrt(norm);
```



Example: Vector norm *reduction*

$$\sqrt{\sum_i v(i) * v(i)}$$


```
norm = 0.0;

#pragma omp parallel default(none) \
    shared( vect ) reduction( + : norm ) private(i)
{
    #pragma omp for
    for (i = 0; i < vleng; i++)
        norm += vect[i]*vect[i];

}

norm = sqrt(norm);
```

// master copy
// lNorm gone!
// priv. copy: 0
// private copy
// comb. copies
// master copy



Example: Vector norm *reduction, parallel for*

$$\sqrt{\sum_i v(i) * v(i)}$$

```
norm = 0.0; // master copy

#pragma omp parallel for default(none) \
    shared( vect ) reduction( + : norm )

    for (i = 0; i < vleng; i++)
        norm += vect[i]*vect[i]; // private copy

norm = sqrt(norm); // master copy
```



Supported operators and initial values for reduction in Fortran (OpenMP 3.0)

Name	Symbol	Initial Value of local copy
add	+	0
multiply	*	1
subtract	-	0
logical AND	.and.	.true.
logical OR	.or.	.false.
EQUIVALENCE	.eqv.	.true.
NON-EQUIV.	.neqv.	.false.
maximum	max	smallest representable number
minimum	min	largest representable number
bitwise AND	iand	All bits on
bitwise OR	ior	0
bitwise XOR	ieor	0



Supported operators and initial values for reduction variables in C (OpenMP 3.0)

Name	Symbol	Initial Value of local copy
add	+	0
multiply	*	1
subtract	-	0
bitwise AND	&	~ 0
bitwise OR		0
bitwise XOR	^	0
logical AND	&&	1
logical OR		0



Restrictions and comments on reduction

- Arrays are unsupported as reduction variables in C/C++
- No pointer or reference types in C/C++
- Fortran ALLOCATABLE must be allocated at the beginning of construct and must not be de-allocated during construct
- No Fortran pointers or assumed size arrays
- No order of threads is specified
 - Repeated runs are typically not bit-identically (common issue in parallel computing)
 - This is a *race condition*, which is typically tolerated!
- OpenMP 4.0: declare your own reductions



User defined reductions

- Allows to define you own reduction operations
- Particularly useful with derived data types, examples:
 - C/C++: struct
 - Fortran: type
- You need:
 - Combiner: combines thread private results to final result
 - Initialiser: initialise private contributions at outset



A case study:

Maximum value and its position

- Problem:
 - We have a large array
 - Determine the maximum value
 - Location of the maximum in the array
- Parallelisation
 - Assign portion of array to each thread
 - Thread determines maximum and position in its part
 - User defined reduction to determine final result



Example in Fortran

- Consider the type:

```
type :: mx_s
    real value
    integer index
end type
```

- Declare a reduction operator named maxloc:

```
!$omp declare reduction(maxloc: mx_s: &
!$omp  mx_combine(omp_out, omp_in) ) &
!$omp  initializer(mx_init(omp_priv, omp_orig))
```

- The operation can be triggered by the name “maxloc”
- Utilises subroutine `mx_combine` and `mx_init`
- Acts on object of type: `mx_s`



The initialiser

- Subroutine or assignment statement – here: subroutine
- Acts on variables:
 - omp_priv: reference to variable to be initialised
 - omp_orig: reference to original variable prior to construct
- Example: Initialise from value prior to construct:

```
subroutine mx_init(priv, orig)
  type(mx_s), intent(out) :: priv
  type(mx_s), intent(in)  :: orig
  priv%value = orig%value
  priv%index = orig%index
end subroutine mx_init
```



The combiner

- Subroutine or assignment statement – here: subroutine
- Acts on variables:
 - omp_in: reference to contribution from thread
 - omp_out: reference to combined result
- Example: replace if contribution is larger

```
subroutine mx_combine(out, in)
  type(mx_s), intent(inout) :: out
  type(mx_s), intent(in) :: in
  if ( out%value < in%value ) then
    out%value = in%value
    out%index = in%index
  endif
end subroutine mx_combine
```



How to use it:

- You can use it similar to predefined reductions:

```
mx%value = val(1)
mx%index = 1

!$omp parallel do reduction(maxloc: mx)
do i=2, count
    if (mx%value < val(i)) then
        mx%value = val(i)
        mx%index = i
    endif
enddo
```

- Easily readable code
- Similar to what one would do in serial programming



Example in C

- Consider the type:

```
struct mx_s {  
    float value;  
    int index;  
};
```

- Declare a reduction operator named maxloc:

```
# pragma omp declare reduction(maxloc: \  
    struct mx_s: mx_combine(&omp_out, &omp_in)) \  
    initializer(mx_init(&omp_priv, &omp_orig))
```

- The operation can be triggered by the name “maxloc”
- Utilises subroutine `mx_combine` and `mx_init`
- Acts on object of type: `mx_s`



The initialiser in C

- Expression – here: implemented with a function
- Acts on variables:
 - omp_priv: reference to variable to be initialised
 - omp_orig: reference to original variable prior to construct
- Example: Initialise from value prior to construct:

```
void mx_init(struct mx_s *priv, struct mx_s *orig)
{
    priv->value = orig->value;
    priv->index = orig->index;
}
```



The combiner in C

- Expression – here implemented with a function
- Acts on variables:
 - omp_in: reference to contribution from thread
 - omp_out: reference to combined result
- Example: replace if contribution is larger

```
void mx_combine(struct mx_s *out, struct mx_s *in)
{
    if ( out->value < in->value ) {
        out->value = in->value;
        out->index = in->index;
    }
}
```



How to use it in C:

- You can use it similar to predefined reductions:

```
mx->value = val[0];  
mx->index = 0;  
  
#pragma omp parallel for reduction(maxloc: mx)  
for (i=1; i < count; i++) {  
    if (mx.value < val[i])  
    {  
        mx.value = val[i];  
        mx.index = i;  
    }  
}
```

- Easily readable code
- Similar to what one would do in serial programming



Declaring a reduction operation

Syntax summary

- Basic syntax in C

```
#pragma omp declare reduction (reduction-identifier : \  
    typename-list : combiner) [initializer-clause] new-line
```

- Basic syntax in Fortran

```
!$omp declare reduction(reduction-identifier :      &  
!$omp type-list : combiner) [initializer-clause]
```



Dealing with global storage

- By default global storage is **shared**
- C/C++ examples for global storage:
 - file scope variables
 - static variable
- Fortran examples for global storage:
 - `COMMON` blocks
 - `module data`
 - variable with `save` attribute
- Not always what is needed



Directive: threadprivate in C

- Directive `threadprivate`
- Each thread gets private copy
- Outside `parallel`: modify copy of master
- Example prints:
 - 4 on master thread
 - 1 else

```
int g_var=1;
#pragma omp \
    threadprivate(g_var)

int main{
    g_var = 4;
#pragma omp parallel
    {
        printf("%d\n", g_var);
    }
return 0;
}
```



Directive: threadprivate in Fortran

- Directive `threadprivate`
- Each thread gets private copy
- Outside `parallel`: modify copy of master
- Example prints:
 - 4 on master thread
 - 1 else

```
module gmod
  integer g_var=1
  !$omp threadprivate(g_var)
end module gmod
```

Program example

```
use gmod
g_var = 4
!$omp parallel
  print *,g_var
!$omp end parallel
End program example
```



Clause: copyin to initialise threadprivate in C

- Initialise threadprivate data from master: copyin clause
- Example prints:
 - 4 on all threads

```
int g_var=1;
#pragma omp \
    threadprivate(g_var)

int main{
    g_var = 4;
#pragma omp parallel \
    copyin(g_var)
{
    printf("%d\n", g_var);
}
return 0; }
```



Directive: threadprivate in Fortran

- Initialise threadprivate data from master: `copyin` clause
- Example prints:
 - 4 on all threads

```
module gmod
  int g_var=1
  !$omp threadprivate(g_var)
end module gmod

Program example
  use gmod
  g_var = 4
  !$omp parallel copyin(g_var)
    print *,g_var
  !$omp end parallel
End program example
```



More on threadprivate

- `threadprivate` data **unchanged** between parallel regions if:
 - Neither region nested inside other parallel region
 - Both regions have same thread count
 - Internal variable *dyn-var* is false in both regions (use function `omp_set_dynamic` to set)
- In Fortran one can make a `COMMON` block `threadprivate`

```
integer :: a, b, c  
COMMON/abccom/a,b,c  
!$OMP threadprivate(/abccom/)
```



Summary

- Discussed special private variables
 - `firstprivate`: initialisation of private variables
 - `lastprivate`: set value of private variable to value of last loop iteration or last section at end of construct
 - `reduction`: Calculating sums, products etc. in parallel
 - `threadprivate`: privatise global storage
- **Remark:** The above will handle standard situations
 - Constructs of the earlier examples for special cases
 - Initialise `private` variable from `shared` variable
 - Atomic/critical writes of shared variables

