



**LUND**  
UNIVERSITY



**HPC2N, UmU**

# Thread binding for OpenMP

Joachim Hein

LUNARC & Centre of Mathematical Sciences

Lund University

Pedro Ojeda May

HPC2N

Umeå University



# OS Thread-scheduling


- Operating system schedules processes & threads
  - Running vs sleeping
  - Assignment to physical hardware
- Processes and threads typically scheduled freely
  - OS moves long running processes with time
  - Utilise as many cores as possible at all time
  - Good: web- or transactionserver, Interactive multiuser
- Can create problems for parallel HPC applications
  - Loss of memory locality
  - Multiple compute threads on same core



# Structure of 16-way AMD Interlagos node

## Example for a CC-NUMA architecture

- 2 Processors
- 4 dual module groups
- Memory attached to these groups (CC-NUMA architecture)
- Hypertransport
- 16 cores



# Thread binding

- Threads allocate memory locally if available
- Thread binding
  - Bind threads to hyper-threads, cores or groups thereof
  - Avoid multiple threads on a single core/hyper-thread
  - Threads are not moved away from “their” memory
  - Allows optimisation of inner node communication
- Up to OpenMP 3.1, thread binding was compiler dependent
  - OpenMP 4.0 provides standardised binding features
- Watch out for hyper-threading



# Mapping of threads on the hardware

- Mapping of threads onto hardware affects performance
- Effect to consider
  - Shared caches facilitating fast data exchange
  - Multiple memory busses / memory controllers
  - Shared functional units (e.g. FPU in AMD Interlagos)
- Particular important for under populated nodes



# OpenMP binding a two stage process

You need to define

- Places
  - Where should the threads run
  - Examples: Hyper-thread, Core, Socket (Processor)
  - Groups of the above
- Affinity policy
  - How threads get distributed on the “Places”



# OpenMP Places

- Controlled by environment variable: OMP\_PLACES
  - Not changeable at runtime
- Defines hardware resources to execute OpenMP threads
  - Can execute more than one OpenMP thread
- Easiest to use abstract names, choices:

<code>export OMP_PLACES=threads</code>	(Hyper-thread)
<code>export OMP_PLACES=cores</code>	(physical Core)
<code>export OMP_PLACES=sockets</code>	(Processor)



# Specifying an OpenMP place

- Specify a place as a list of positive numbers
  - Labeling smallest execution unit exposed by environment
  - Typically hardware thread
- Specifying places
  - You can define a place as comma separated list in {}
    - E.g: {0, 4, 8, 12} – system can choose from 4 units
  - Define a place as interval with a colon:
    - E.g: {4:5} – this corresponds to {4, 5, 6, 7, 8}
    - The 5 give the number of units, not the end!
  - You can define a stride:
    - E.g: {5:4:3} – this corresponds to {5, 8, 11, 14}





# The OMP\_PLACES variable as a list

- Comma separated list of places

```
export OMP_PLACES="{0,1,2,3},{4:4},{8:4}"
```

- You can specify (strided) intervals for the place list

```
OMP_PLACES="{0,1,2,3}:3:5"
```

corresponds to:

```
OMP_PLACES="{0,1,2,3},{5,6,7,8},{10,11,12,13}"
```



# Thread binding: Affinity policy

- Affinity policy is controlled by the environment variable `OMP_PROC_BIND`

<code>OMP_PROC_BIND</code>	effect
<b>false</b>	no binding
<b>true</b>	binding enabled
<b>master</b>	bind to same place as master
<b>close</b>	bind subsequent threads to subsequent places, round-robin when all used
<b>spread</b>	bind threads as far away as possible

- You can provide a list containing `master`, `close` and `spread`, that will be used for nested parallel regions



# Declaring binding on the parallel region

## OpenMP 4.0

- One can declare binding in your source for a parallel region
- Binding (OMP\_PROC\_BIND) must not be false
- Fortran:

```
!$omp parallel proc_bind(policy)
```

- C

```
#pragma omp parallel proc_bind(policy)
```

- Valid values for *policy* are:

master

close

spread



# Effect of Binding

## Calulating a vector norm

- icc 17.0.1
- Haswell node:  
2 Intel E5-2650 v3
- Benchmark:
  - Initialising the vector and calculating the norm
  - Separate loops, cache inefficient
- Small problem: Communication bound
  - Staying inside a processor most efficient
- Large problem: Streams bound
  - Utilising both processors more efficient

Treads	Size	None	Close	Spread
10	4000	4.7 $\mu$ s	2.5 $\mu$ s	4.7 $\mu$ s
20	4000	6.2 $\mu$ s	5.0 $\mu$ s	5.1 $\mu$ s
10	400M	96ms	179ms	94ms
20	400M	90ms	90ms	90ms



# Summary

- Explained the need for thread binding
- Since OpenMP 4.0: binding part of the OpenMP standard
  - Previously compiler dependent
- Demonstrated effect of binding on a simple benchmark

