


The Task directive

Joachim Hein

LUNARC & Centre of Mathematics
Lund University




Lund University 1

1




There is more than regular loops

- So far discussed: regular structures
 - Loops with known start and end
 - Fortran array constructs
- Many problems have irregular structures:
 - Recursive
 - Linked lists
 - Loops with unknown end (e.g. while)
 - Divide and Conquer
 - ...
- Depending on the details they might still be **parallelisable**




Lund University 2

2

THE TASK CONSTRUCT




Lund University 3

3

The task directive in Fortran

- creates “explicit task” from
 - code body
 - data environment at that point
- place inside parallel region
- execution:
 - now or later
 - encountering or other thread

```
!$omp task [clauses]
  code body
!$omp end task
```



Lund University 4

4

The task directive in C

- creates “explicit task” from
 - code body
 - data environment at that point
- place inside parallel region
- execution:
 - now or later
 - encountering or other thread

```
#pragma omp task [clauses]
{
    code body
}
```



Allowed data sharing attributes for tasks

- `private`
 - data is private to the task
- `firstprivate`
 - data is private to the task
 - data initialised when task directive is encountered
- `shared`
 - data is shared – **only** way to return a result!
- `default`
 - Fortran: `shared` | `private` | `firstprivate` | `none`
 - C: `shared` | `none`



Data sharing without a default

- When no default is declared on a “task” directive:
- If shared by **all** implicit tasks in the current team:

Variable is: `shared`

- Otherwise

Variable is: `firstprivate`

- **My recommendation:** `default(none)`

Lund University

7



7

Example for execution: Thread encountering the following

- Executes “code block 1”
- Creates a task for “code block 2”
- May
 - Execute the task for “code block 2”
 - Pick up another task
 - Continue with “code block 3”
- At some point:
 - Has to execute code block 3
- No control
 - Who executes code block 2
 - When code block 2 is finished


```
code block 1

!$omp task
  code block 2
!$omp end task

code block 3
```

Lund University

8



8

Controlling when a task finishes

```
!$omp taskwait
```

- Ensures **child tasks** have completed
- Does not consider grand-children etc.

```
!$omp barrier
```

- Ensures all tasks in the innermost parallel region have finished
- Remark: Instead of waiting, thread can execute task generated elsewhere



Allowing suspension of the current task

- At a taskyield construct
 - Can suspend the current task to:
 - Execute a different task

Fortran:

```
!$omp taskyield
```

C:

```
#pragma omp taskyield
```




Using taskgroup to control tasks finishing (OpenMP 4.0)

- A taskgroup construct, defines a region with an implied task scheduling point at the end
- Current task suspended until **all** descendant tasks (incl. grand children etc.) completed

```
!$omp taskgroup
  do i=1, n
    !$omp task ...
      call processing(...)
    !$omp end task
  end do
!$omp end taskgroup
```

This waits for all task,
incl. tasks generated
in processing



Lund University

11


11

Using taskgroup to control tasks finishing (OpenMP 4.0)

- A taskgroup construct, defines a region with an implied task scheduling point at the end
- Current task suspended until **all** descendant tasks (incl. grand children etc.) completed

```
#pragma omp taskgroup
{ for (int i=0; i<n; i++)
{
  #pragma omp task ...
  { processing(...);
  }
}
}
```

This waits for all task,
incl. tasks generated
in processing



Lund University

12

12

Controlling task creation

- Creating a task encounters significant overhead
 - Requires significant work inside to pay off
- Use `if`-clause to control task creation


```
!$OMP task if(level .lt. 10) ...
      ...
!$OMP end task
```
- In case expression evaluates to `.false.`
 - Encountering thread executes code body directly (included task)



Lund University

13

13

Final tasks

- A task can carry a `final` clause


```
!OMP task final( level .gt. 30) ...
      ...
!OMP end task
```
- If expression evaluates to `.true.`, all encountered tasks will be:
 - included
 - final



Lund University

14

14

Mergeable tasks

- A task can be declared as mergeable

```
!$omp task mergeable ...
```

```
#pragma omp task mergeable ...
```




- In case of an undeferred or included task, the implementation may:
 - Use the data environment of the generating task (incl. internal control variables)
 - Use for optimisation, e.g. with `final`



Task scheduling points


- Threads may switch to different task at **task scheduling point**
- Task scheduling points are:
 - Immediately after generation of explicit task
 - After point of completion of a task
 - At `taskwait`, `taskyield`
 - At `barrier` (explicit or implicit)
 - At the end of `taskgroup`
- Untied tasks (not in course) may switch at any point
 - Care with `critical`, locks etc.
 - E.g. task may switch out of critical region
→ deadlock!





1st Case study

PARALLELISING A RECURSIVE ALGORITHM




Lund University 17

17

Fibonacci numbers

- Mathematical series:
 $F_0 = 0$
 $F_1 = 1$
 $F_n = F_{n-1} + F_{n-2}$
- First numbers in series:
0, 1, 1, 2, 3, 5, 8, 13, 21, 34, ...
- Recursive program: not efficient!!!



Lund University 18

18

Fortran code: Serial version

```

recursive function recursive_fib(in) result(fibnum)
  integer, intent(in) :: in
  integer(lint) :: fibnum, sub1, sub2
  if ( in .gt. 1) then

      sub1 = recursive_fib( in - 1 )

      sub2 = recursive_fib( in - 2 )

      fibnum = sub1 + sub2
  else
      fibnum = in
  endif
end function recursive_fib

```

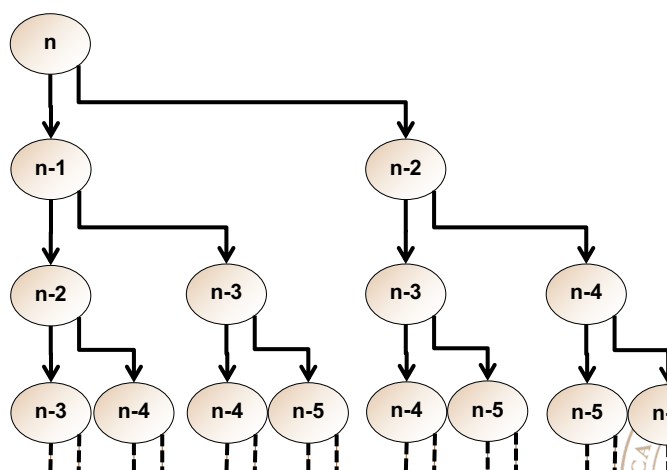
Lund University

19




19

Recursive Fibonacci



Lund University

20



20

Fortran code: Serial version

```
recursive function recursive_fib(in) result(fibnum)
  integer, intent(in) :: in
  integer(lint) :: fibnum, sub1, sub2
  if ( in .gt. 1) then

    sub1 = recursive_fib( in - 1 )

    sub2 = recursive_fib( in - 2 )

    fibnum = sub1 + sub2
  else
    fibnum = in
  endif
end function recursive_fib
```

Lund University

21



21

Fortran code: Towards a parallel version I

```
recursive function recursive_fib(in) result(fibnum)
  integer, intent(in) :: in
  integer(lint) :: fibnum, sub1, sub2
  If ( in .gt. 1) then
    !$OMP task shared(sub1) firstprivate(in)
    sub1 = recursive_fib( in - 1 )
    !$OMP end task
    sub2 = recursive_fib( in - 2 )

    fibnum = sub1 + sub2
  else
    fibnum = in
  endif
end function recursive_fib
```

sub1 shared
declared inside function

in firstprivate
initilised at task creation

Lund University

22



22

Fortran code: Towards a parallel version II

```

recursive function recursive_fib(in) result(fibnum)
  integer, intent(in) :: in
  integer(lint) :: fibnum, sub1, sub2
  If ( in .gt. 1) then
    !$OMP task shared(sub1) firstprivate(in)
      sub1 = recursive_fib( in - 1 )
    !$OMP end task
    !$OMP task shared(sub2) firstprivate(in)
      sub2 = recursive_fib( in - 2 )
    !$OMP end task

    fibnum = sub1 + sub2
  else
    fibnum = in
  endif
end function recursive_fib

```

place task at 2nd call

Problem:
Need to have sub1 & sub2

Lund University

23

23

Fortran code: Parallel version

```

recursive function recursive_fib(in) result(fibnum)
  integer, intent(in) :: in
  integer(lint) :: fibnum, sub1, sub2
  If ( in .gt. 1) then
    !$OMP task shared(sub1) firstprivate(in)
      sub1 = recursive_fib( in - 1 )
    !$OMP end task
    !$OMP task shared(sub2) firstprivate(in)
      sub2 = recursive_fib( in - 2 )
    !$OMP end task
    !$OMP taskwait
    fibnum = sub1 + sub2
  else
    fibnum = in
  endif
end function recursive_fib

```

Solved: taskwait
Waits for the 2 tasks above
Recursion takes care of grand-children

Lund University

24

24

How to call?

The original serial code

```

Program fibonacci
  !$ use omp_lib
  integer, parameter :: lint = selected_int_kind(10)
  integer(lint) :: fibres
  integer :: input
  read (*,*) input

  fibres = recursive_fib(input)

  print *, "Fibonacci number", input, " is:", fibres
End program fibonacci

```

Lund University

25



25

How to call?

Need to start a parallel region

```

Program fibonacci
  !$ use omp_lib
  integer, parameter :: lint = selected_int_kind(10)
  integer(lint) :: fibres
  integer :: input
  read (*,*) input
  !$OMP parallel shared(input, fibres) default(none)

  fibres = recursive_fib(input)


  !$OMP end parallel
  print *, "Fibonacci number", input, " is:", fibres
End program fibonacci

```

Problem:
Each thread starts
Fibonacci calculation :o

Lund University

26



26

How to call?

```

Program fibonacci
  !$ use omp_lib
  integer, parameter :: lint = selected_int_kind(10)
  integer(lint) :: fibres
  integer :: input
  read (*,*) input
  !$OMP parallel shared(input, fibres) default(none)
    !$OMP single
      fibres = recursive_fib(input)
    !$OMP end single
  !$OMP end parallel
  print *, "Fibonacci number", input, " is:", fibres
End program fibonacci

```

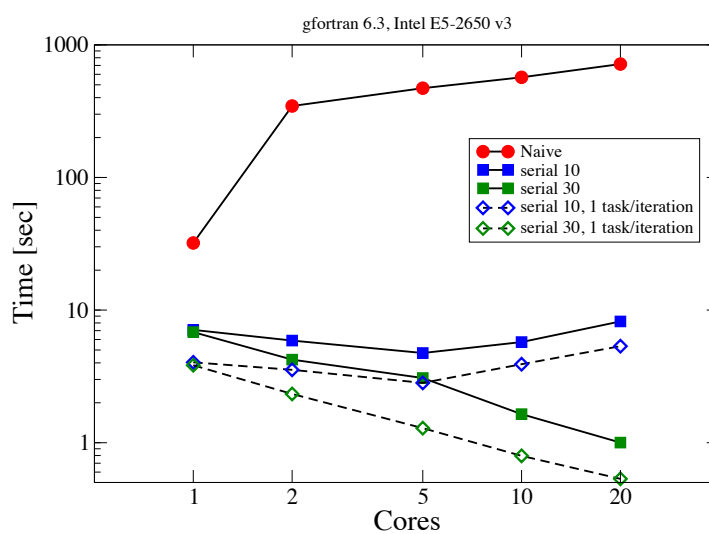
Lund University

27



27

Performance for Fibonacci number 40




Lund University

29



29



30




Discussion of Fibonacci performance

- Naïve implementation
 - 2 tasks per iteration
 - poor performance
- Using if-clause
 - No tasks created for low input value: helps
- Start only 1 task/iteration
 - Helps more
- Too little work per task
- Limit number of tasks
- Hardware:
 - 2 socket/server
 - Intel E5-2650 v3
 - 10 cores per Processor
- gfortran
 - Version 6.3
 - Thread binding
- Intel ifort
 - Version 17.1
 - Thread binding

Lund University

31


31

HPC2N, UmU

2nd Casestudy

SELF-REFINING RECURSIVE INTEGRATOR




Lund University 32

32

Mesh refinement

- Codes endloying irregular grids
- Dynamic grid refinement / coarsening allows efficiency
 - Example: Fluiddynamic
 - refine grid were eddy develops
 - coarsen when eddy vanishes
- Case study: self refining integrator for 1D function



Lund University 33

33

Basic algorithm

- Evaluate function at 5 regular space points in interval
- Estimate integral:
 - Polygon using all 5 points
 - Polygon using only 3 points (first, centre, last)
- Check difference between the two integrals
 - Compare to threshold times interval length
- If accurate add contribution to accumulation variable
- If not accurate:
 - Split interval into two pieces
 - Run integrator on both pieces (Recursion)



Lund University

34

34

Implementation details: Parallel region

- Use shared variable `accumulator` for result
 - Declared as module variable
- Start recursive integrator from `single`
- The implied barrier ensures tasks are finished here
- Recursive subroutine: `rec_eval_shared_update`

```
accumulator = 0.0D0
!$OMP parallel default(none) &
!$OMP shared(accumulator) &
!$OMP shared(startv, stopv, &
!$OMP unit_err, gen_num)

!$OMP single
call rec_eval_shared_update( &
startv, stopv, unit_err, gen_num)
!$OMP end single
!$OMP end parallel
```

Lund University

35

35

Implementation details: Task startup (in: rec_eval_shared_update)

```
!$OMP task shared(accumulator) firstprivate(my_start, my_stop) &
!$OMP default(none) firstprivate(my_gen, u_err) &
!$OMP if(task_start)
  call rec_eval_shared_update( &
    my_start, 0.5_dpr * (my_start + my_stop), u_err, my_gen)
!$OMP end task

!$OMP task shared(accumulator) firstprivate(my_start, my_stop) &
!$OMP default(none) firstprivate(my_gen, u_err) &
!$OMP if(task_start)
  call rec_eval_shared_update( &
    0.5_dpr * (my_start + my_stop), my_stop, u_err, my_gen)
!$OMP end task
```

Lund University

36

36

Implementation details: Result accumulation


- Declare shared variable (global to avoid issue in ifort 19u5):
atomic update when accurate
- Declare threadprivate variables
 - Thread executing task achieving accuracy adds to his threadprivate copy
 - After barrier (implied in end single) atomic update of threadprivate data into shared variable
- **Remarks:**
 - Careful: threadprivate and task scheduling points
 - Value can be changed after scheduling point
 - threadprivate isn't private to the task
 - OpenMP 5.0 has reduction constructs for tasks

Lund University

37

37

Testfunction: $\sin^2(10000 x) \sin^4(x)$




Lund University

38

38

Where does it sample




Lund University

39

39

Performance results



- Task started every 5th generation
- Poor results with atomic updates (there are millions)
- threadprivate accumulation satisfactory result


Lund University

40



40

Performance results



- Task started every 5th generation
- Poor results with atomic updates (there are millions)
- threadprivate accumulation satisfactory result
- Efficient utilization of up to 128 cores

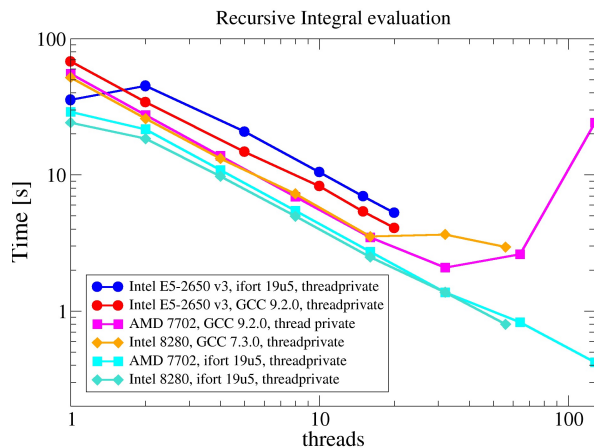
Lund University

41



41

Comparing compilers



- GCC: inferior scalability beyond 20 cores

Lund University

42



42

Task dependencies (OpenMP 4.0)

- One can declare dependencies


```
!$omp task depend (type : list)
```

```
#pragma omp task depend (type : list)
```

- 3 dependency types:
 - in: task depends on all previous siblings with an out or inout dependency on one or more of the list items
 - out, inout: task depends on all previous siblings with an in, out or inout dependency on one or more of the list items
- The list are variables, that may include array sections

Lund University

43



43

Example for task dependency

- Code snippet:

```
#pragma omp task depend (out: a)
    task_function_1( &a );
#pragma omp task depend (in: a)
    task_function_2( a );
#pragma omp task depend (in: a)
    task_function_3( a );
```

- Wait for task_function_1 to be finished
- Can execute task_function_2 and task_function_3 in any order on any thread

Lund University

44



44

Tasks and reductions (OpenMP 5.0)

- A taskgroup around the reduction area is required
- Reduction variable included in a task_reduction clause of the taskgroup:

```
task_reduction( operator : variable_list )
```

- Reduction variable include in a in_reduction clause of the task construct

```
in_reduction( operator : variable_list
```

Lund University

45



45

Example in C

```
double sum = 0.0;
#pragma omp single
{
    #pragma omp taskgroup task_reduction(+: sum)
    {
        for (int i = 0; i < N; i++)
            #pragma omp task in_reduction(+: sum)
            {
                sum += evaluation(i);
            }
    }
}
```

Lund University

46



46

Taskloop (OpenMP 4.5)

- The taskloop construct distributes loops on tasks
 - Similar to the loop construct
- By default the taskloop implies a taskgroup

Lund University

47



47

Basic syntax

```
!$OMP taskloop default(none) shared(...) private(...)
do i = 1, N
    ...
enddo
```

```
#pragma omp taskloop default(none) shared(...) private(...)
for (i=0; i<N; i++)
{
    ...
}
```

Lund University

48

48

Clauses for taskloop

- Clauses introduced before
if(*scalar-expr*), shared, private,
firstprivate, lastprivate, default, collapse,
final(*scalar-expr*)
- A reduction clause was introduced for the taskloop
was introduced in OpenMP 5.0
- An in_reduction clause is available if the taskloop is
embedded in a taskgroup with a task_reduction
- The clause nogroup removes the implied taskgroup
- There is a taskloop simd construct

Lund University

49

49

Controlling the number of task created

- Use `grainsize` or `num_tasks` to control task number
 - Only one allowed
- `grainsize` controls number of loop iterations per task
 - each tasks gets between `grainsize` and `2*grainsize` iterations
- `num_tasks` specifies the number of tasks created
- Additional restrictions from iteration count



Summary

- Explained the task construct
 - Task scheduling
 - Task completion
- Discussed performance aspects in case studies
 - Need to control the number
 - Decent amount of work per task

