

Vectorisation

JOACHIM HEIN, LUNARC & CENTRE OF MATHEMATICAL SCIENCES

PEDRO OJEDA MAY, HPC2N, UMEÅ UNIVERSITY



HPC2N, UmU



Overview

- Vector registers
- SIMD construct
- Declare SIMD construct to vectorise functions





LUND
UNIVERSITY

Vectorisation



HPC2N, UmU



Modern hardware has wide registers

Overview on x86 system

Instruction set	Register width	Single prec. words	Double prec. words	Typical hardware
SSE, SSE2	128 bit	4	2	modern x86
AVX, AVX2	256 bit	8	4	x86 since 2011
AVX-512	512 bit	16	8	Skylake Knights Landing

- Concept also exists in non-x86 hardware, examples:
 - ARM: NEON
 - IBM Power: VSX



Example:

AVX2 FMA instruction

- AVX: 256 bit registers - 4 doubles
- Single instruction - 8 flops:


$$a_1 + b_1 * c_1$$

$$a_2 + b_2 * c_2$$

$$a_3 + b_3 * c_3$$

$$a_4 + b_4 * c_4$$

- Enable via compiler option:
 - Without cross compilation:
 - » GCC: `-march=native -O3`
 - » Intel: `-xHost -O3`
 - Cross compilation: explicit specification



Basic example for SIMD deployment

```
do i=1, n  
    a(i) = b(i) + c(i)  
enddo
```

- Execute multiple loop iterations simultaneously
- Reduce loop count accordingly
- Iterations need to be independent



What needs to be done for SIMD (Simplified)

```
do i=1, n, 4
  a(i  ) = b(i  ) + c(i  )
  a(i+1) = b(i+1) + c(i+1)
  a(i+2) = b(i+2) + c(i+2)
  a(i+3) = b(i+3) + c(i+3)
enddo
```

- Execute multiple loop iterations simultaneously
- Iterations need to be independent
- Compiler might need to add a peel



Basic example for SIMD deployment

```
for (i=1; i<n; i++)  
{  
    a[i] = b[i] + c[i]  
}
```

- Execute multiple loop iterations simultaneously
- Reduce loop count accordingly
- Iterations need to be independent



What needs to be done for SIMD (Simplified)

```
for (i=0; i<n; i+=4)
{
    a[i  ] = b[i  ] + c[i  ];
    a[i+1] = b[i+1] + c[i+1];
    a[i+2] = b[i+2] + c[i+2];
    a[i+3] = b[i+3] + c[i+3];
}
```

- Execute multiple loop iterations simultaneously
- Reduce loop count accordingly
- Compiler might need to add a peel



Automatic vectorisation

- Modern compilers vectorise many loops automatically
 - Choose right instruction set and optimisation level
 - » GNU: `-O3 -march=native`
 - » Intel: `-O3 -xHost`
 - Compilers can report on vectorisation
 - » GNU: `-fopt-info-vec -fopt-info-vec-missed`
 - » Intel: `-qopt-report -qopt-report-phase=vec`
- Compiler needs help in complex situations
 - OpenMP SIMD construct: portable way to help





LUND
UNIVERSITY

SIMD construct



HPC2N, UmU



simd construct (Fortran)

- Assure compiler that the following loop can be vectorised

```
!$omp simd  
do i = 1, n  
    a(i) = a(i) + b(i)  
enddo
```



simd construct in C

- Assure compiler that the following loop can be vectorised

```
#pragma omp simd
for (i = 1; i < n; i++)
{
    a[i] = a[i] + b[i];
}
```



Clauses for `simd` construct

- Data sharing:
`private`, `lastprivate`, `reduction`
- There is no default(`none`) here!
- Number of loops associated with construct
`collapse(n)`



Clauses for `simd` construct (cont.)

- Clause `safelen` allows vectorisation of certain dependencies

```
!$omp simd safelen(7)
do i = 1, n
    a(i) = a(i) + a(i+7)
enddo
```

- Allowed to load up to 7 values in the register
- This would be difficult to parallelise



Clauses for `simd` construct (cont.)

- Clause `safelen` allows vectorisation of certain dependencies

```
#pragma omp simd safelen(7)
for (i = 1; i < n; i++)
    a[i] = a[i] + a[i+7];
```

- Allowed to load up to 7 values in the register
- This would be difficult to parallelise



Clauses `simdlen`

- Clause `simdlen`: preferred number of consecutive iterations:

```
!$omp simd simdlen(4)
do i = 1, n
    a(i) = a(i) + b(i)
enddo
```

- This will suggests to do 4 iterations simultaneously



Clauses `simdlen`

- Clause `simdlen`: preferred number of consecutive iterations:

```
#pragma omp simd simdlen(4)
for (i = 1; i < n; i++)
{
    a[i] = a[i] + b[i];
}
```

- This suggests to do 4 iterations simultaneously



Clause linear

- Declare a linear relationship between iteration (\neq loop index) and a variable

```
j=0
!$omp simd linear(j:2)
do i = 1, N, 3
    j = j + 2
    a(i) = b(j)
enddo
```

- Data sharing clause, j is now private



Clause linear

- Declare a linear relationship between iteration (\neq loop index) and a variable


```
j=0;
#pragma omp simd linear(j:2)
for (i = 1; i < n; i++)
{
    j = j + 2;
    a[i] = b[j];
}
```

- Data sharing clause, j is now private




SIMD and data alignment

- Aligned load:



- Un-aligned load (typically not supported)



SIMD and data alignment

- Byte alignment of data
 - An address is e.g. 32-byte-aligned if:
byte-address is divisible by 32
- SIMD loads/stores typically have alignment requirements
 - Unaligned data needs either
 - » Peel loops – load individual number until boundary
 - » Crossload
 - Checking requires extra code – performance impact



Aligning data

- There is no function in OpenMP ☹

- In C one might want to lib-functions from stdlib:

```
int posix_memalign(void **memptr, size_t alignment, size_t size);  
void *aligned_alloc( size_t alignment, size_t size );
```

- In Fortran

- Wrap posix_memalign or aligned_alloc

- Compiler specific tools:

- » Example: Intel compiler

```
ifort -align array32byte
```

aligns all arrays at 32-byte boundaries



Advanced optimisation: aligned

- If you understand/control your data alignment
- Declare it to the system (e.g. 32 byte alignment)

```
!$omp simd aligned(x,32)  
Do i = 1, N  
    x(i) = 2.0D0 * x(i)  
enddo
```

- If correct, this will reduce overheads (peeled loops)
- If false, illegal instructions
- No optional parameter – implementation default alignment



Advanced optimisation: aligned

- If you understand/control your data alignment
- Declare it to the system (e.g. 32 byte alignment)

```
#pragma omp simd aligned(x,32)
for (i = 1; i<N; i++)
    x[i] = 2.0D0 * x[i];
```

- If correct, this will reduce overheads (peeled loops)
- If false, illegal instructions
- No optional parameter – implementation default alignment





LUND
UNIVERSITY

Vectorisation of functions and subroutine



HPC2N, UmU



Function/subroutine calls in loop

- Try to avoid calls in loop, due to performance impact
- Use `declare simd` to create vector versions of functions and subroutines

```
function addfunc(a,b)
  !$omp declare simd(addfunc)
    implicit none
    double precision :: a, b, addfunc
    addfunc = a + b
end function addfunc
```



Using a simd-ised function

- The function can be used in a simd loop

```
!$omp simd  
do i = 1, N  
    c(i) = addfunc(a(i), b(i))  
enddo
```



Function/subroutine calls in loop

- Try to avoid calls in loop, due to performance impact
- Use `declare simd` to create vector versions
- Required in the header **and** source file!

```
#pragma omp declare simd  
double addfunc(double a, double b) {  
    double r;  
    r = a + b;  
    return r;  
}
```



Using a simd-ised function

- The function can be used in a simd loop

```
#pragma omp simd  
for (i = 1; i < N; i++) {  
    c[i] = addfunc(a[i], b[i])  
}
```



Clauses for declare simd

- Discussed before:
 - `simdlen(length)`
 - `linear(linear-list[: linear-step])`
 - `aligned(arg-list[: alignment])`
- Specific clauses
 - `uniform(argument-list)`
value invariant for all invocations
 - `inbranch`
always called inside a conditional statement
 - `notinbranch`
never called inside a conditional statement



Example 1

```
function vecop(a, b, i, offset)
!$omp declare simd(vecop) uniform(a,b,offset) &
!$omp linear(i:1)
    integer :: i
    double precision :: a(*), b(*), offset, vecop
    vecop = a(i) + b(i) + offset
end function vecop
```

- GCC 6.3, 8.3, 9.2 and clang 9.0.1 will not compile this!
 - Complain about a and b in uniform
 - Possible workaround: scalar code
- Intel does compile this



Using function vecop

```
double precision :: a(N), b(N)

...

!$omp simd
Do i = 1, N
    a(i) = vecop(a, b, i, 3.1d0)
enddo
```



Example 2

```
function cube(x)
!$omp declare simd inbranch
    implicit none
    double precision :: cube, x
    cube = x*x*x
end function cube
```

- This will generate code with a mask and suppress the unmasked version



Using the cube function

```
double precision :: x(N), y(N)
...
!$omp simd simdlen(4)
do i = 1, N
    if( y(i).gt.0.0d0 ) then
        y(i) = cube(x(i))
    endif
enddo
```

- Will operate on 4 long vectors, but not do the operation on lane, where condition is false (applying masking)

Combined constructs

- Distribute a loop and simd-ise it:

```
!$omp do simd
```

- Start parallel region, distribute a loop and simd-ise it:

```
!$omp parallel do simd
```

- Modified schedule

```
!$omp do simd schedule(simd:static, 11)
```

– new chunksize = (chunksize/simdlen) * simdlen

– always a multiple of the simdlen



Summary

- Using the SIMD construct to assist in loop vectorisation
- declare SIMD construct to allow vectorised function calls
- Some compilers allow for SIMD part of OpenMP only
 - Intel: `-qopenmp-simd`

